

Query Optimization

Cost-Based

SWEN 304
Trimester 2, 2017

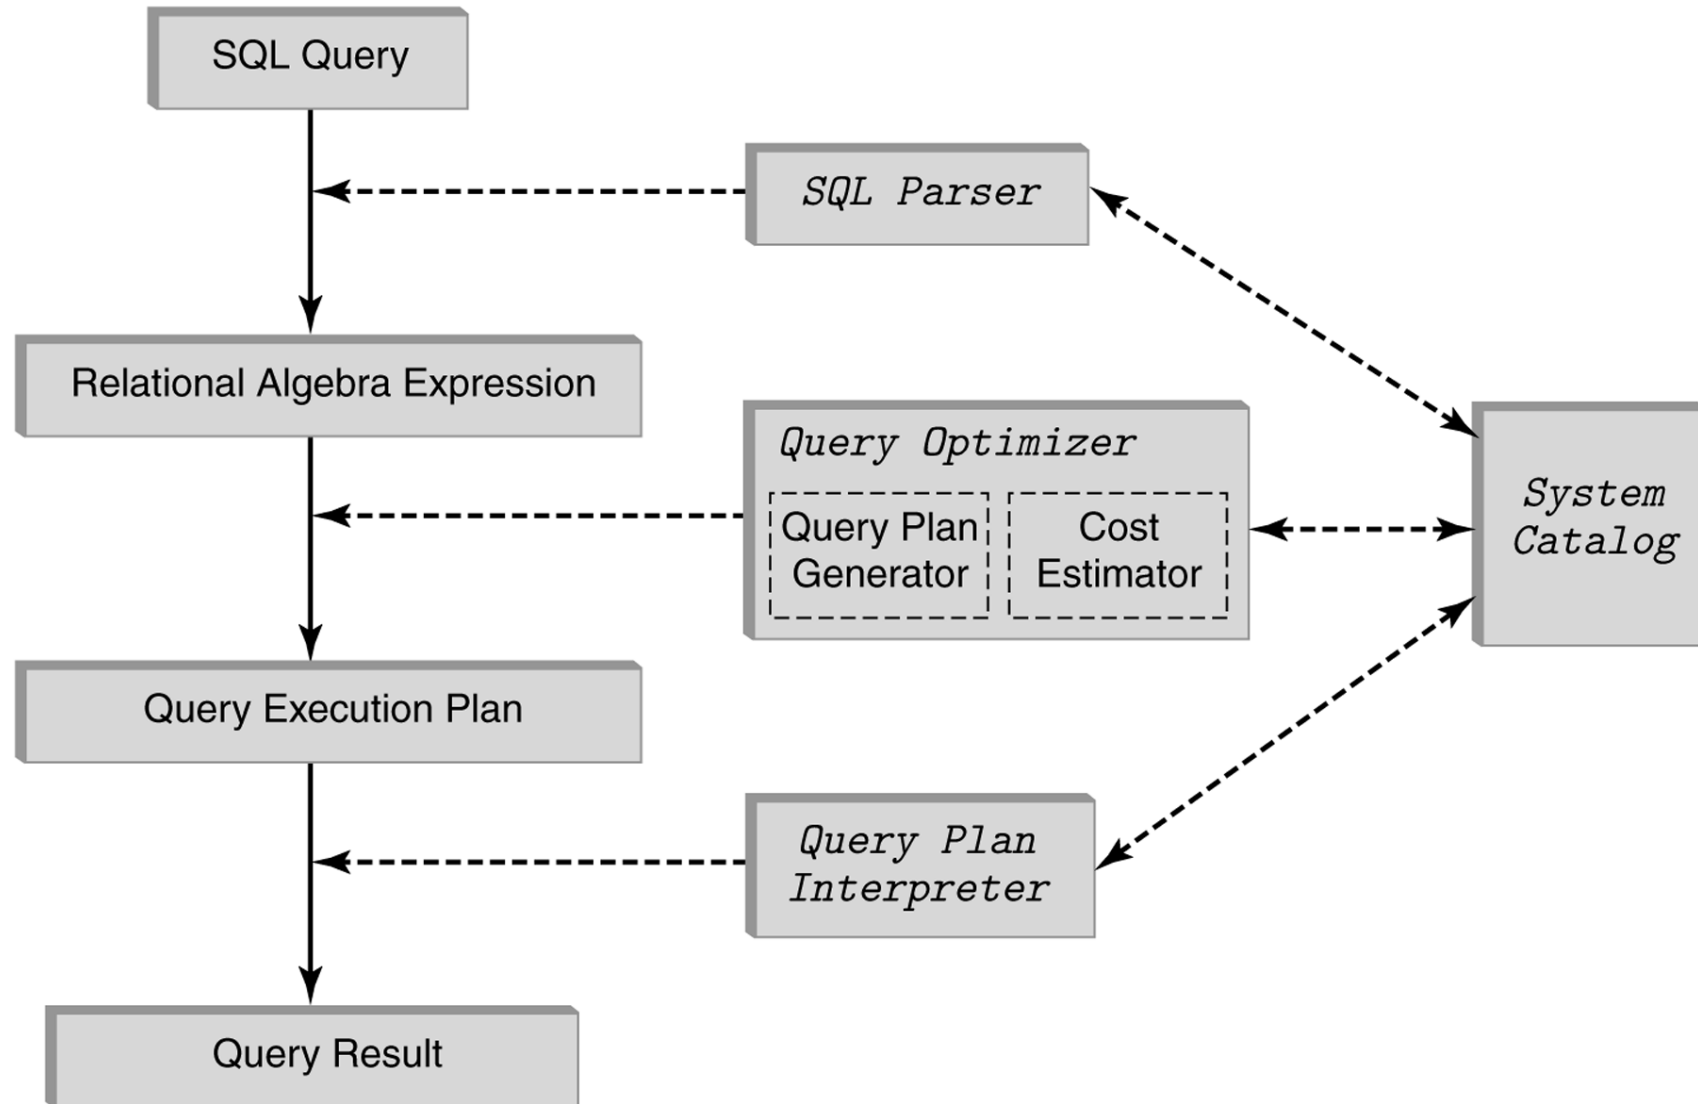
Lecturer: Dr Hui Ma

Engineering and Computer Science



Slides by: Pavle Morga & Hui Ma

Outline – Query Evaluation in DBMS



Outline – Cost-based Query Optimization

- Why cost-based optimization?
- How to measure cost of query operations?
 - Cost function of a projection, selection, join, ...
- How to evaluate query operations?
 - Algorithms for projection, selection, join, ...
- Query tree of physical operators
 - Reading: chapters 17, 18, 19 of 6/E of the textbook
 - Supposed knowledge **File Organization**

Example

R (Reserves)

<u>sid</u>	<u>bid</u>	<u>day</u>
58	101	10/09/13
22	103	11/09/13
31	103	11/09/13

S (Sailor)

<u>sid</u>	sname	rating	age
22	Jerry	7	25
31	Tom	8	30
58	Minny	10	22

B (Boats)


<u>bid</u>	<u>bname</u>	<u>color</u>
101	Dragon	blue
102	Moon	red
103	Star	green
104	Clipper	blue
105	Mary	green

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

Example

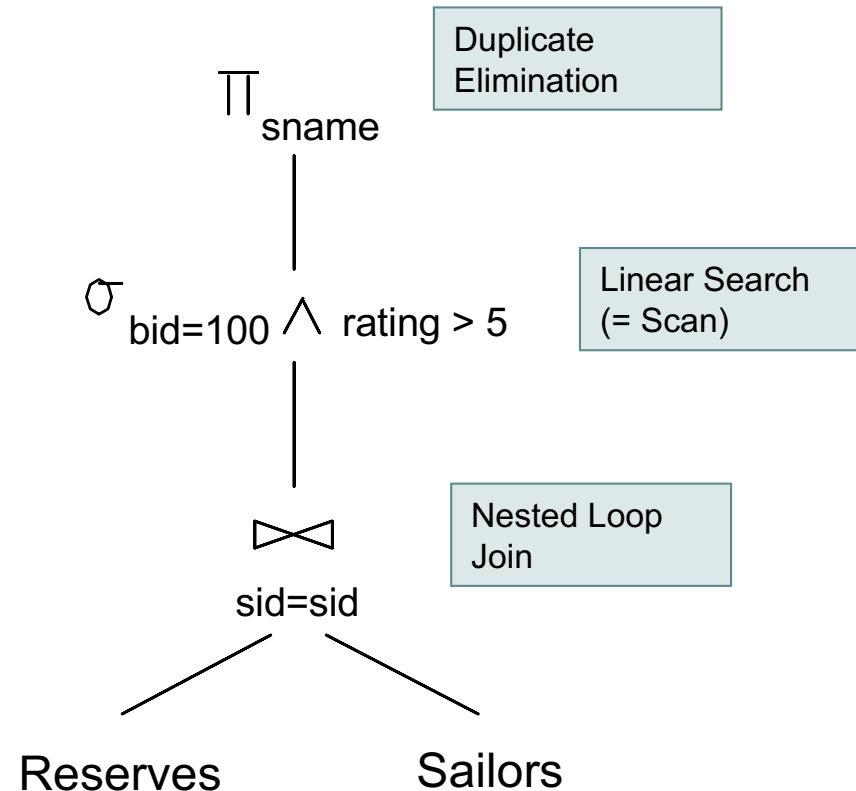
- Here is a query tree for the SQL query
 - Is this optimal?
 - How good is it?
 - How to measure “goodness”?
 - Need to estimate its costs!

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```



Example

- Here is an execution plan:
 - How to evaluate the operations?
 - Need to pick an evaluation algorithm for each operation
 - DBMS typically offer multiple evaluation algorithms
 - that have been implemented by the DBMS programmers



Cost-Based Optimization

- A good query optimizer does not rely solely on heuristic rules
- It chooses that query execution plan which has the **lowest cost estimate**
 - Or at least one whose costs are **reasonably good**
- After heuristic rules are applied to a query, there still remain a number of alternative ways to execute it
 - These alternative ways rely on the existence of **different auxiliary data structures** and **algorithms**
- Query optimizer estimates the cost of executing each of alternative ways, and chooses the one with the **lowest cost**

Cost Components of a Query Execution

- Secondary storage **access** cost:
 - Reading data blocks during data searching,
 - Writing data blocks on disk, and
 - Storage cost (cost of storing intermediate files)
- Computation cost (CPU cost)
- Main memory cost (buffer cost)
- Communication cost
- Very often, only secondary storage access cost is considered
- So, the cost **C** will be the number of **disk accesses**

Cost Related Catalog Content


- For the purpose of a query cost estimate, a Catalog should contain following information for each **base relation**:
 - Number of tuples (= records) r
 - Number of blocks b
 - Blocking factor f (= the number of tuples that fit into one block)
 - Available access methods and access attributes:
 - Access methods: sequential, indexed, hashed
 - Access attributes: primary key, indexing attributes, sort key
 - The number of levels h of each index
 - The number of distinct values d of each attribute

Some Assumptions

- To make it simpler, we shall suppose that:
 - All tuple fields are of a **fixed size** (although variable field size tuples are very frequent in practice)
 - All the intermediate query results are **materialized** (although there are some advanced optimizers that apply pipelined approach)
 - **Materialized**: intermediate query results are **stored on a disk** as temporary relations
 - **Pipelining**: tuples of the intermediate results are subjected to all subsequent operations **without temporary storing**
 - Intermediate results produced by unary operations are retained in blocks of the size as the initial files

Example

- **Materialized evaluation:**
 - evaluate one operation at a time, starting from the bottom
 - intermediate results are materialized into temporary relations (write to disk)
 - E.g., the result of the join is materialized, the temporary relation is then read from disk to compute the selection
 - Similarly, the result of the selection is materialized, and then from disk to compute the projection



Example

- Materialized evaluation is always applicable
 - Cost of writing results to disk and reading them back can be quite high
- Pipelined evaluation:
 - evaluate several operations together, passing the results of one operation on to the next
 - E.g., don't store result of join, but pass them on to selection
 - Similarly, don't store result of selection, but pass them on to projection
 - Much cheaper than materialization: no need to store a temporary relation to disk

Cost Function of a Project Operation

- Suppose the `<attribute_list>` of a project operation
 - contains a relation schema **key**, or
 - the keyword DISTINCT is **not** used in the SQL SELECT command,
- Then project operation
 - Reads all b_1 blocks, containing r tuples of the size n from the secondary storage into memory, and
 - Writes back b_2 blocks, containing r tuples of the size m ($m < n$), on the secondary storage
- Since $m < n$, it follows $b_2 < b_1$
 - The cost function is
$$C = b_1 + b_2$$
- Complexity $O(r)$

Projection: Evaluation of DISTINCT

- Suppose the <attribute_list> of the SELECT clause does **not** contain a relation schema **key** and the keyword **DISTINCT** is used
- Algorithm:
 - Read b_1 blocks from the base relation,
 - Drop unwanted columns and write b_2 blocks with duplicate tuples back
 - Look for duplicate tuples by reading in each of b_2 , and comparing its tuples with tuples in the other blocks,
 - Number of reads:
$$\sum_{i=2}^{b_2} i = \frac{b_2^2 + b_2 - 2}{2}$$
 - Complexity $O(r^2)$

Projection: Reduce Cost of DISTINCT with Sorting

- Algorithm:
 - Read b_1 blocks from the base relation,
 - Drop unwanted columns and write b_2 blocks with duplicate tuples back
 - **Sort** b_2 blocks (Complexity $O(r * \log r)$)
 - Read successive blocks from the sorted file and **drop** all duplicates (except one)
 - Write back b_3 blocks without duplicates
- Overall complexity $O(r * \log r)$

Projection: Reduce Cost of DISTINCT with Index

- For SELECT DISTINCT <attr_list> a costly sort can be avoided if there exists an appropriate secondary index on the whole <attr_list>
- Then, the query optimizer only has to traverse index, to retrieve all different secondary key values, and write them back as an intermediate file
- So, let $d(Y)$ ($\leq r$) be the number of different $Y =$ <attr_list> values, and m the number of node entries, and suppose index is implemented as a B^+ -tree
- The total cost will be:
$$C = \lceil d(Y) / m \rceil + \lceil d(Y) / f \rceil$$
- Complexity: $O(d)$ or $O(r)$

Selection: Attribute Selection Cardinality

- If an attribute A of a relation schema N has $d(A)$ actual distinct values, then its **selection cardinality** $s(A)$ is

$$s(A) = r / d(A)$$

- For a key K , $d(K) = r$, and $s(K) = 1$
- If an attribute A is not a key, then

$$s(A) = (r / d(A)) \geq 1$$

- Selection cardinality $s(A)$ of the attribute A , allows us to compute how many tuples is expected to contain a given value

$$a \in \pi_A(N)$$

- **We always assume a uniform distribution**

Cost Function of a Select Operation

- Let s ($0 \leq s \leq r$) be the number of tuples that satisfy selection condition
- In the case of a **linear** search, select operation is performed by
 - reading b blocks in, and
 - writing s tuples as $\lceil s / f \rceil$ blocks back
- The cost function is
$$C = b + \lceil s / f \rceil$$
- Complexity $O(r)$

Cost Functions of Join Operation

- The join operation is one of the most time consuming operations in query processing
- We shall consider joins $N \bowtie_{jc} M$, where N and M are relational variables, and jc is a join condition of the form $N.Y = M.Y$
- N is called outer loop relation, and M is called inner loop relation
- Of the four basic join algorithms, we consider in more detail only **nested-loop** join

The Size of a Join Result

- We start from a simplifying supposition that the join condition $jc \equiv N.Y = M.Y$ is based on a **foreign key/primary key pair**, and that the corresponding referential integrity constraint $M[Y] \subseteq N[Y]$ is satisfied

- Then and only then, the size of a join result is

$$|N \bowtie_{jc} M| \leq r_M$$

- The number of blocks is

$$\lceil r_M / f \rceil$$

where f is the blocking factor

Question for You

- If Y is the primary key of N , and the referential integrity constraint $M[Y] \subseteq N[Y]$ is satisfied, then

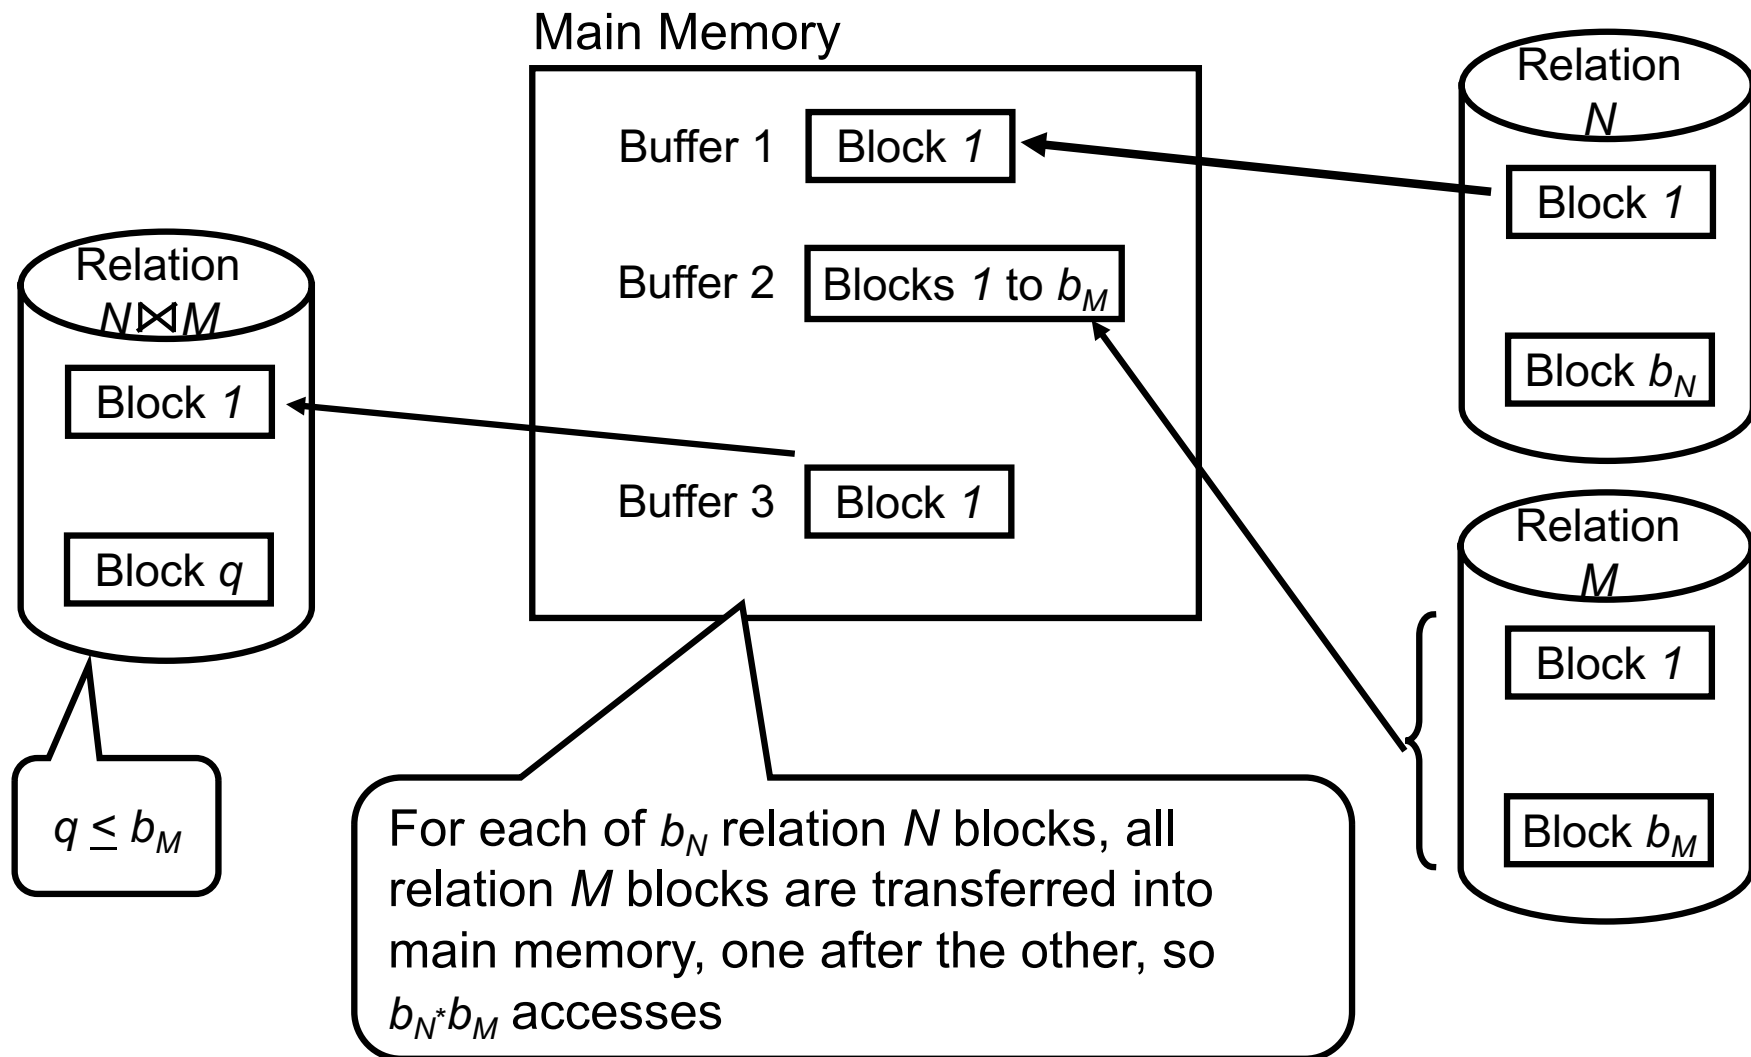
$$|N \bowtie_{jc} M| \leq r_M$$

- Why is the statement true?

Nested-Loop Join

- Algorithm:
 - For each tuple t in N (outer loop relation),
 - retrieve each tuple u from M (inner loop relation),
 - test whether the two tuples satisfy join condition jc
(whether $t[Y] = u[Y]$)
 - If yes, concatenate t and u and include them into the join result
- Let n ($n > 2$) be the number of buffers available for storing:
 - $n - 2$ out of b_N outer relation blocks at once
 - 1 of b_M inner relation blocks, and
 - 1 of $\lceil r_M / f \rceil$ result blocksin the main memory

Nested-Loop Join – Three Buffers



Cost of Nested-Loop Join

- The cost of nested-loop join:

$$C = b_N + b_M \lceil b_N / (n - 2) \rceil + \lceil r_M / f \rceil$$

With total number of blocks read for outer relation is b_N and total number of blocks read for inner relation is $b_M \lceil b_N / (n - 2) \rceil$

- The number of **buffers** n has considerable impact on the number of disk accesses C
- Since

$$b_M \lceil b_N / (n - 2) \rceil \approx b_N \lceil b_M / (n - 2) \rceil$$

the number of disk accesses C will be smaller if

$$b_N < b_M$$

- Complexity $O(r^2)$

Other Join Algorithms

- Single-loop join
 - Uses index on inner relation
 - Only exceptionally better than the others
- Sort-merge join
 - Relies on sorting,
 - A reasonably good choice
- Hash join
 - Relies on hashing
 - Greedy on memory
 - If provided enough memory, the best choice

Single-Loop Join (Homework)

- The prerequisite:
 - An index (or hash key) on join attribute $M.Y$ of the inner relation M
- Algorithm:
 - Retrieve each tuple of the outer loop relation N and use the access structure to retrieve directly all matching tuples of the inner loop relation M
- Complexity $O(r * \log r)$ (large constants neglected)
- Needs at least **four buffers**, but additional buffers bring only a slight improvement
- Can't apply after select and project

Sort-Merge Join (Homework)

- Algorithm:
 - Sort the N and M relations ($m + 1$ buffers needed, $m \geq 2$)
 - Read successive and sorted blocks of N and M into memory
 - Compare successive N and M tuples from the blocks read in
 - Include the tuples that match into join result
- Complexity $O(r * \log r)$
 - Which is the complexity of sorting
- Frequently the best choice

Hash Join (Homework)

- Consider a function h that will be applied on the join attributes $N.Y$ and $M.Y$
- Both relations N and M are split (one after the other) into m partitions using the same hash function h
- That way, partitions N_i and M_i contain tuples that are equivalent with regard to h , and a tuple from N_i may join only with some tuples from M_i
- Pairs (N_i, M_i) of partitions are joined one after the other and stored into join result (m iterations needed)
- Complexity $O(r)$

Question for You

- We introduced four join algorithms:
 - Nested-Loop Join (exhaustive search),
 - Single-Loop Join (B-tree, not applicable after select or project),
 - Sort-Merge Join (sort is expensive),
 - Hash Join (high demand on memory buffers)
- How does a query processor find the most effective one?
 - a) Randomly takes one
 - b) Calculates the cost of using each of them and takes the least expensive

Efficiency Improvement Techniques

- The efficiency of executing relational algebra operations can be improved using:
 - Indexes (select, project with DISTINCT),
 - More memory (join, sort)
 - Sorting:
 - project with DISTINCT,
 - set theoretic operations,
 - aggregates with GROUP BY
- More on this in the tutorial

Combining the Optimization Techniques

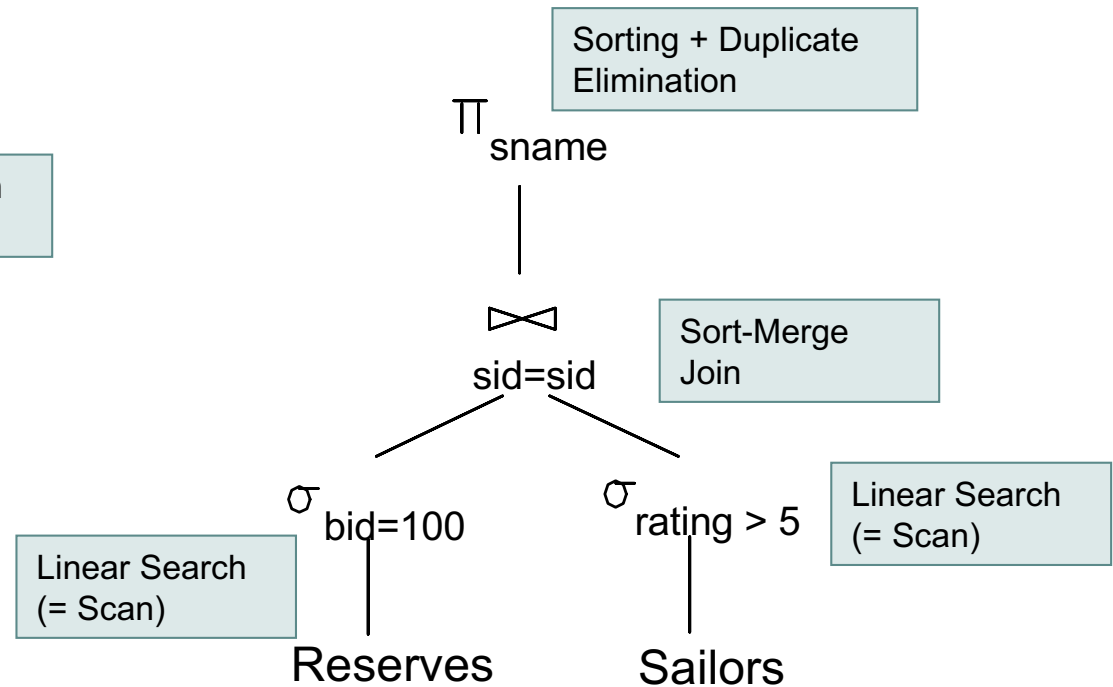

- The cost based optimization is applied onto the query tree that is a result of the heuristic optimization
- Cost based optimization mainly examines:
 - Implementation methods of operations, and
 - The order of multiple joins and their implementation methods
- Query tree of physical operators is produced when the **cheapest** execution plan is provided with access methods and algorithms to be used in executing the relational algebra operations

Nested Query Optimization

- Optimization of a nested query depends on whether it is correlated or not:
 - In a non correlated nested query, the result of the inner SELECT has to be computed only once, and each tuple of the outer SELECT is compared to that result
 - In a correlated nested query, inner SELECT is evaluated for each tuple of the outer SELECT
 - Namely, result of the inner SELECT depends on the attribute values of the current outer SELECT tuple

Example

- Which execution plan is better?



Summary

- DBMS processes a declarative query by converting it to the query tree of logical operators, and by optimizing it
 - looks for a reasonably efficient strategy to implement a query
- Heuristic optimization and cost based optimization are two basic optimization techniques
- Cost based optimization is applied to the result of heuristic optimization
 - exhaustive analyze the number of disk accesses of alternative available methods and algorithms to execute a query
- Techniques that improve query cost:
 - Indexing,
 - Using larger memory,
 - Sorting