

# Java DataBase Connectivity

SWEN 304

Trimester 2, 2017

Lecturer: Dr Hui Ma

**Engineering and Computer Science**



Slides by: Pavle Mogin & Hui Ma

# Plan for Java DataBase Connectivity (JDBC)

- Motivation
- Architecture
- JDBC Classes and Interfaces
  - JDBC Driver Management
  - Controlling transaction behavior
- Executing SQL statements
- Obtaining result data
- Matching data types
- Exceptions
- Closing a connection
  
- *Further Reading: The Java™ Tutorial:*  
`www.mcs.vuw.ac.nz/technical/java/tutorial/index.html`

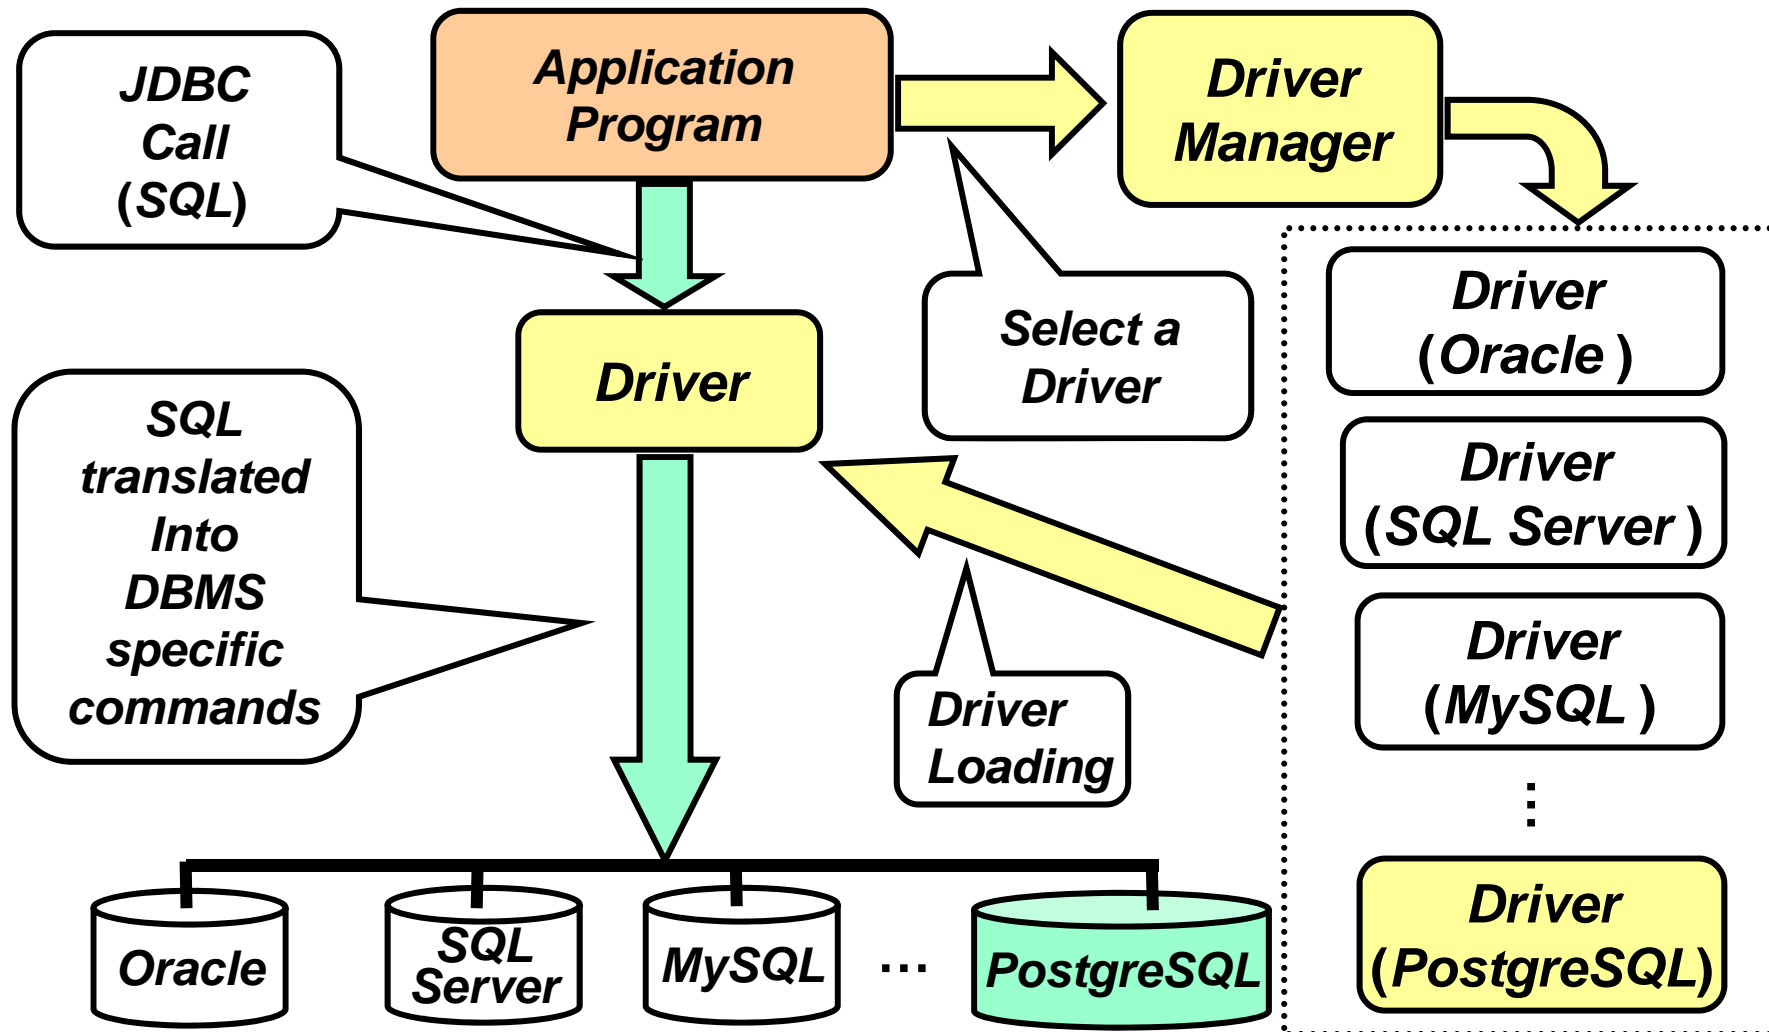
# Motivation for Using JDBC

- In practice, databases are not only accessed by human users through the user interface, but also by application programs
- Application programs are written in a general purpose programming language (GPPL)
  - for example, Java, C/C++, Python, ...
- When developing applications, software engineers use SQL for data management from inside their application program
- The application program is written in a GPPL with SQL statement **embedded** into it
  - therefore, the GPPL is also called the **'host' language**

# The Java Database Connectivity (JDBC) API

- **Java Database Connectivity (JDBC)** is the standard application program interface (API) for accessing databases from a Java program
  - it allows us to embed SQL statements into Java code
- JDBC is supported by all relevant DBMS, including all major commercial and open source DBMS
  - Interactions of an application program with a specific DBMS are accomplished through a DBMS-specific **JDBC driver**
- Application programs using JDBC are independent of the particular DBMS that is used
  - Independence holds on source code & executable code level
  - Independence is a huge advantage when developing applications for multiple DBMS
  - Independence is achieved by an **extra level** of indirection

# Accessing Databases from Application Programs



# JDBC Architecture (1)

- When using JDBC the following play a role :
  - Application programs
  - A database system (DBMS plus databases)
  - A DBMS-specific JDBC driver
  - A driver manager

## JDBC Architecture (2)

- Application programs ...
  - Dynamically load the JDBC **drivers** needed,
  - Initiate a **connection** with a database,
  - Set transaction **boundaries** (BEGIN,..., {COMMIT | ROLLBACK}),
  - Acquire **locks**,
  - Submit **SQL** statements,
  - **Receive** data,
  - **Process** data,
  - **Process** error messages,
  - Decide whether to **commit** or **roll-back** a transaction,
  - Disconnect from the database to **terminate** a session

## JDBC Architecture (3)

- The DBMS ...
  - Processes data manipulation commands,
  - Returns results to the application program
  
- The JDBC driver ...
  - Establishes **connection** with a database,
  - Submits data manipulation **requests**,
  - Accepts **results** returned by the DBMS,
  - Translates DBMS specific **data types** into Java data types,
  - Translates **error** messages
  
- The driver manager ...
  - Loads and supervises available JDBC drivers for various DBMS
  - Calls JDBC drivers to connect to a database



# JDBC Classes

- JDBC is a collection of Java classes and interfaces
  - All these are put together in the `java.sql` package
- It contains methods for:
  - **Connecting** to a remote database,
  - **Executing** SQL statements,
  - **Iterating** over a set of tuples from a SQL statement,
  - **Transaction** management,
  - **Exception** handling

# JDBC DriverManager class

- JDBC provides the `DriverManager` class
- Among others, it defines methods to enable dynamic addition and deletion of JDBC drivers :
  - `registerDriver()`
  - `deregisterDriver()`
- The first step in **connecting to a database** (managed by some DBMS) is to load a suitable JDBC driver for that particular DBMS
- Any JDBC 4.0 drivers that are found in the class path are automatically loaded

# Registering a JDBC Driver

- However, drivers prior to JDBC 4.0 must be loaded manually with the method `Class.forName`  
`public static native Class.forName (String name)`  
`throws ClassNotFoundException`
  - This creates a `Driver` object for the respective JDBC driver
- Example: to load a JDBC driver for PostgreSQL use
  - `Class.forName( "org.postgresql.Driver" );`
- If you are unsure which JDBC drivers will be available, load them manually

# Establishing a Connection

- To connect to a database in a DBMS, the `getConnection` method of the `DriverManager` object is called
  - ```
Connection con =  
    DriverManager.getConnection(url, [userId],  
    [password]);
```
  - This method requires a URL to the database
- In the application program, the method call starts a **session** with the database by creating a `Connection` object

# Specifying the URL for the Connection

- The URL `url` to the database is of the form  
`jdbc:[drivertype]:[database]`
- Herein,
  - `jdbc` is a constant,
  - `[drivertype]` is the type of the database we want to connect (e.g. `postgresql`), and
  - `[database]` is the address of the database in form  
`//hostname[:portnumber]/database_name`
- Example:
  - `jdbc:postgresql://db.ecs.vuw.ac.nz/007_jdbc`

# Connection Interface

- The interface `java.sql.Connection` has a number of classes and methods that are used:
  - To control **transactional** behavior of a `Connection` object,
  - To create and execute **SQL** statements,
  - To iterate over the **result** returned by a DBMS, and
  - To finish interaction with a database by **closing** the connection
- After acquiring a connection (with the name say `con`) and before it is closed, the same connection can be used for executing several transactions

# Controlling Transaction Behavior-Start

- By default, a Connection **automatically** commits changes after executing each SQL statement
- The method  
`public abstract void setAutoCommit(  
boolean autoCommit) throws SQLException`  
is applied onto a Connection object
- To designate the **start** of a transaction (BEGIN point), we assign a value `false` to `autoCommit`  
`con.setAutoCommit( false );`

# Controlling Transaction Behavior-End

- A transaction is **terminated** using:
  - Either  
`public abstract void commit() throws  
SQLException`  
or  
`public abstract void rollback() throws  
SQLException`
- And (after any of them)  
`con.setAutoCommit(true)`  
on the `Connection` object



# What Next?

- Executing SQL statements
  - Statement object
  - PreparedStatement object
- Obtaining result data
- Matching data types
- Exceptions
- Closing a connection

# Executing SQL Statements

- JDBC supports three different ways of executing SQL statements:
  - **Statement,**
  - **PreparedStatement, and**
  - `CallableStatement`

## Statement Class and It's Subclasses

- The **Statement** class is the base class of the three classes used to submit queries to a DBMS
  - Its objects are used to send such SQL queries to a DBMS that are executed with no repetition within a transaction and that have no parameters
- **PreparedStatement** objects are used for SQL statements with parameters or for those that are executed multiple times (in a loop)
  - SQL statements of PreparedStatement objects may be precompiled yielding better performance
- **CallableStatement** objects are used with stored procedures and are out of the scope of the course

# Submitting a SQL Query to a DBMS

- The following steps should be performed in order to submit a SQL statement to a DBMS either using a Statement (**S**) or PreparedStatement (**PS**) object:

1. Define a SQL query as a string
2. Create a S or PS object

If the SQL statement is one of CREATE, INSERT, DELETE, UPDATE, or SET type:

3. Apply **executeUpdate()** method onto the S or PS object

Else (the SQL statement is of the **SELECT** type):

3. Create a **ResultSet** object
4. Feed into the **ResultSet** object the return value of applying **executeQuery()** method onto the S or PS object

## Statement Objects With executeUpdate

```
String insert="INSERT INTO Grades " +  
"VALUES (007007,'C305','A+')";
```

```
Statement stmt=con.createStatement();  
int return_value =  
stmt.executeUpdate(insert);
```

- For INSERT, DELETE, and UPDATE queries, the return value will be the **number** of tuples affected
- For CREATE or SET, the return value should be **0**

# ResultSet Object

- The `executeQuery()` method returns an object of the type **set** (or **superset**)
- This set object should be assigned to an object of the **ResultSet** class
- The `ResultSet` class has the `next()` method that allows traversing the set in a tuple at a time fashion
- Initially, the `ResultSet` object is positioned **before** the first tuple of the result
- The method `next()` returns `true` if there is a **next** tuple in the result, otherwise `false`
- After executing `next()`, the `ResultSet` object contains a pointer to the current tuple

## Statement Objects With `executeQuery`

```
String select="SELECT * FROM Grades"  
+ "WHERE StudentId=007007";
```

```
Statement stmt =  
con.createStatement();
```

```
ResultSet rs =  
stmt.executeQuery(select);  
while (rs.next()) {  
    // extracting data from rs tuples  
    // data processing  
}
```

## Extracting Data from the Result

- To match Java and database data types, JDBC specifies mappings and provides accessor methods in the `ResultSet` class

```
int j_studId;  
String j_courseId;  
String j_grade;  
while (rs.next()) {  
    j_studId=rs.getInt("StudentId");  
    j_courseId=rs.getString("CourseId");  
    j_grade=rs.getString(3)  
    // 3 is the column number in the result  
}
```



## PreparedStatement With executeUpdate

```
String insert="INSERT INTO Grades VALUES (?, ?, ?)";
PreparedStatement prstmt =
    con.prepareStatement(insert);
boolean end=false;
while(!end){
    // suppose j_studId, j_courseId, j_grade, and end
    are
    // dynamically initialized to desired values
    prstmt.setInt(1, j_studId);
    prstmt.setString(2, j_courseId);
    prstmt.setString(3, j_grade);
    int return_value = prstmt.executeUpdate();
    ...
}
```

## PreparedStatement with executeQuery

```
String select= "SELECT * FROM Grades  
WHERE StudentId = ?";
```

```
PreparedStatement prstmt =  
    con.prepareStatement(select);
```

```
// suppose j_studId is initialized on the  
// desired value
```

```
prstmt.setInt(1, j_studId);
```

```
ResultSet rs =  
prstmt.executeQuery();
```

```
while(rs.next()) {  
  
}
```

## Closing a Connection

- Before exiting from an application program all connections acquired should be closed by applying

```
public abstract void close() throws  
SQLException
```

method on each of them

# Exceptions

- Most of the methods in `java.sql` can throw an exception of the type `SQLException` if an error occurs
- In addition to inherited `getMessage()` method, `SQLException` class has two additional methods for providing error information:
  - `public String getSQLState()` that returns an SQL state identifier according to SQL:1999 standard, and
  - `public int getErrorCode()` that retrieves a vendor specific error code
- Each JDBC method that throws an exception has to be placed inside a **try** block followed by a **catch** block

## Code to See Exceptions

```
try{
    /* Code that could generate an
       exception goes here. If an
       exception is generated, the
       catch block below will print out
       information about it*/
}
catch (SQLException ex){
    System.println(ex.getMessage());
    System.println(ex.getSQLState());
    System.println(ex.getErrorCode());
}
```

## Summary

- JDBC Transactions are executed by:
  - Acquiring a Driver,
  - Constructing a connection object,
  - Establishing transaction boundaries
  - Submitting SQL statements,
  - Retrieving results,
  - Processing either results returned or exception errors,
  - Committing or roll-backing transactions, and
  - Disconnecting from databases to terminate interaction
- **Statement** objects - SQL statements have no parameters,
- **PreparedStatement** objects - SQL statements have parameters