

Database System Engineering



SWEN304 **SQL tutorial**

Hui Ma, Yi Mei
hui.ma@ecs.vuw.ac.nz
yi.mei@ecs.vuw.ac.nz

Structured Query Language (SQL)

- We are concerned with three major problems
 - How to **define** a database schema and its relation schemas?
 - How to **store and manipulate** data in relations?
 - How to **retrieve** data from a database?
- SQL is used as
 - a Data **Definition** Language (DDL)
 - a Data **Manipulation** Language (DML)
 - a **Query** Language (QL)

SQL as DDL

- Catalog
- Domains
- Tables (Relation)
- Constraints
- Modify schemas (DROP and ALTER)

Catalog

- The catalog records **essential information** (meta data) about this database
 - **Schema** table: information on all relation schemas
 - **Attribute** table: information on all attributes
 - **Data types** table: information on all data types (domains)
 - **Constraints** table: information on all constraints

SQL Definitions

```
CREATE SCHEMA UNIVERSITY AUTHORIZATION huima;
```

```
CREATE DOMAIN idno
```

```
AS INT
```

```
DEFAULT 300001
```

```
NOT NULL
```

```
CONSTRAINT idnoconstr
```

```
CHECK (VALUE > 300000 AND VALUE <= 399999);
```

UNIVERSITY Database

UNIVERSITY = {STUDENT(StudId, Lname, Fname, Major),
COURSE(CourId, Cname, Points, Dept),
GRADES(StudId, CourId, Grade)}

STUDENT			
StudId	Lname	Fname	Major
300111	Smith	Susan	COMP
300121	Bond	James	MATH
300143	Bond	Jenny	MATH
300132	Smith	Susan	COMP

COURSE			
CourId	Cname	Points	Dept
COMP302	DB sys	15	Engineering
COMP301	softEng	20	Engineering
COMP201	Pr & Sys	22	Engineering
MATH214	DisMat	15	Mathematics

GRADES		
StudId	CourId	Grade
300111	COMP302	A+
300111	COMP301	A
300111	MATH214	A
300121	COMP301	B
300132	COMP301	C
300121	COMP302	B+
300143	COMP201	ω
300132	COMP201	ω
300132	COMP302	C+

Relation Schema: COURSE

```
CREATE TABLE COURSE (  
    CourId CHAR(7) CONSTRAINT cspk PRIMARY KEY,  
    CName CHAR(15) NOT NULL,  
    Points INT NOT NULL  
        CONSTRAINT pointschk  
        CHECK (Points >= 0 AND Points <= 50),  
    Dept CHAR(25)  
);
```

With/Without Domain

```
CREATE TABLE STUDENT (  
    StudId INT  
        NOT NULL  
        DEFAULT 30000  
        CONSTRAINT stpk PRIMARY KEY  
        CONSTRAINT StdRange CHECK  
            (StudId BETWEEN 300000 AND 399999),  
    LName CHAR(15) NOT NULL,  
    FName CHAR(15) NOT NULL,  
    Major CHAR(25) DEFAULT 'Comp'  
);
```

```
CREATE TABLE STUDENT (  
    StudId idno CONSTRAINT stpk PRIMARY KEY,  
    LName CHAR(15) NOT NULL,  
    FName CHAR(15) NOT NULL,  
    Major CHAR(25) DEFAULT 'Comp'  
);
```


Second Key

- Unique and Not Null

```
CREATE TABLE STAFF (  
    staff_id INT PRIMARY KEY,  
    ird_number CHAR(7) NOT NULL UNIQUE,  
    address VARCHAR(255)  
);
```

DROP

- **CASCADE** behavior means deleting the construct itself and all the other constructs related to it,
- **RESTRICT** behavior means that the construct will be deleted only if it is empty (schema), or not referenced by any other construct (like: table, attribute, view)
- e.g. DROP TABLE COURSE RESTRICT
 DROP TABLE STUDENTS CASCADE

ALTER

```
ALTER TABLE STUDENT ADD NoOfPoints INT DEFAULT  
320;
```

```
ALTER TABLE GRADES ALTER Grade SET DEFAULT 'C';
```

```
ALTER TABLE GRADES DROP CONSTRAINT gsri ;
```

```
ALTER TABLE STUDENT DROP CONSTRAINT stpk  
CASCADE;
```

SQL as DML

UNIVERSITY = {STUDENT(StudId, Lname, Fname, Major),
 COURSE(CourId, Cname, Points, Dept),
 GRADES(StudId, CourId, Grade)}

$IC = \{ \text{GRADES}[\text{Id}] \subseteq \text{STUDENT}[\text{Id}],$
 $\text{GRADES}[\text{Course_id}] \subseteq \text{COURSE}[\text{Course_id}] \}$

STUDENT			
StudId	Lname	Fname	Major
300111	Smith	Susan	COMP
300121	Bond	James	MATH
300143	Bond	Jenny	MATH
300132	Smith	Susan	COMP

COURSE			
CourId	Cname	Points	Dept
COMP302	DB sys	15	Engineering
COMP301	softEng	20	Engineering
COMP201	Pr & Sys	22	Engineering
MATH214	DisMat	15	Mathematics

GRADES		
StudId	CourId	Grade
300111	COMP302	A+
300111	COMP301	A
300111	MATH214	A
300121	COMP301	B
300132	COMP301	C
300121	COMP302	B+
300143	COMP201	ω
300132	COMP201	ω
300132	COMP302	C+

SQL as DML

- Three commands used to modify the database:
 - INSERT, DELETE, and UPDATE

- Full insert

```
INSERT INTO STUDENT  
VALUES (111111, 'Bole', 'Ann', Math);
```

- Partial insert

```
INSERT INTO STUDENT (FName, LName, StudId )  
VALUES('Ann', 'Bole', 111111);
```

UPDATE and DELETE

```
UPDATE GRADES  
  SET Grade = 'A+'  
 WHERE CourId = 'C302';
```

```
DELETE FROM STUDENT  
 WHERE FName = 'Susan';
```

```
DELETE FROM STUDENT  
 WHERE StudId IN  
   (SELECT s.StudId  
    FROM STUDENT s, GRADES g  
    WHERE s.StudId = g.StudId AND CourId =  
          'C302');
```

DROP vs DELETE

- DELETE statement performs **conditional** based deletion, whereas DROP command deletes **entire** records in the table
- DELETE statement removes only the rows in the table and it **preserves the table structure** as same whereas DROP command **removes all the data in the table and the table structure**
- DELETE operation **can be rolled back** and it is not auto committed, while DROP operation **cannot be rolled back** in any way as it is an auto committed statement
- DROP is a **DDL** statement while DELETE is a **DML** statement

SQL as QL

- **SELECT** is the basic SQL statement for retrieving data from a database

```
SELECT    [ DISTINCT ]    <attribute_list>  
FROM      <table_list>  
[ WHERE   <condition>    ]
```


SQL as QL

- SQL does not automatically eliminate duplicate tuples in query results
- Use the keyword `DISTINCT` in the `SELECT` clause
 - Only distinct tuples should remain in the result

The University Database

- Suppose for each student only pass grades are recorded in the database

STUDENT

LName	FName	<u>StudId</u>	Major
Smith	Susan	131313	Comp
Bond	James	007007	Math
Smith	Susan	555555	Comp
Cecil	John	010101	Math

COURSE

PName	<u>CourId</u>	Points	Dept
DB Sys	C302	15	Comp
SofEng	C301	15	Comp
DisMat	M214	22	Math
Pr&Sys	C201	22	Comp

GRADES

<u>StudId</u>	<u>CourId</u>	Grade
007007	C302	A+
555555	C302	D
007007	C301	A
007007	M214	A+
131313	C201	B-
555555	C201	C
131313	C302	D
007007	C201	A
010101	C201	D

Single Table Queries

- Retrieve the first and last names of Comp students

```
SELECT FName, LName  
FROM STUDENT  
WHERE Major = 'Comp';
```

FName	LName
Susan	Smith
Susan	Smith

- Find all different grades

```
SELECT DISTINCT Grade  
FROM GRADES ;
```

Grade
A+
A
B-
C

Arithmetic Operations, Sorting

- SQL provides capability to perform four basic arithmetic operations (+, -, *, /) that can be applied to numeric attributes and constants only

```
SELECT 2 + 2;
```

- Sorting of the query result tuples is done using `ORDER BY { <attribute_name> [(ASC|DESC)], ... }` clause after the WHERE clause (ASC is default)

```
SELECT *  
FROM GRADES  
ORDER BY StudId ASC, CourId DESC;
```

Qualification and Aliasing

- Attributes in different relation schemas can have the same names. How do we prevent ambiguity?
- In the `SELECT` clause, we prefix attributes by table name: `SELECT STUDENT.StudId ...`
- In the `FROM` clause, specify a tuple variable from the table: `...FROM COURSE c, GRADES g, STUDENT s`
- In the `WHERE` clause, prefix an attribute by the tuple variable: `WHERE c.CourId = g.CourId`

Join

- To retrieve data from more than one table, we need a new operation: JOIN
- There are different joins:
 - INNER (theta join, equi-join, natural join)
 - OUTER (left, right, full)
 - Most often, we use the equi-join

A JOIN Example

```
SELECT * FROM r1, r2 WHERE r1.B = r2.B ;
```

Join
Condition

r1

<u>A</u>	B
1	b1
2	b1
3	b2
4	b3
5	ω

r2

<u>B</u>	C
b1	c1
b2	c1
b3	c1
b4	c2

r1 EQUI-JOIN r2

<u>A</u>	B	B	C
1	b1	b1	c1
2	b1	b1	c1
3	b2	b2	c1
4	b3	b3	c1

- If there is no join condition what is the result?

Explicit Join

- Q1: Retrieve first name, course id and corresponding grades of the student with Student Id = 007007

```
SELECT FName, CourId, Grade  
FROM (STUDENT NATURAL JOIN GRADES )  
WHERE StudId = 007007 ;
```

- The FROM clause contains a single joined table

Types of Join

- Inner Joins:
 - JOIN, INNER JOIN, EQUI-JOIN, NATURAL JOIN,
- Outer Joins:
 - LEFT OUTER JOIN, RIGHT OUTER JOIN, FULL OUTER JOIN
 - The keyword OUTER may be omitted
 - CROSS JOIN is used to specify the CARTESIAN PRODUCT operation and should be used only with the utmost care

JOIN

- Using JOIN operator

```
SELECT * FROM R1 JOIN R2 ON R1.B = R2.B;
```

R1	
A	B
a ₁	w
a ₂	b ₁

R2	
B	C
b ₁	c ₁
b ₂	c ₁
b ₃	c ₁

R1 JOIN R2			
A	B	B	C
a ₂	b ₁	b ₁	c ₁

- Using NATURAL JOIN operator

```
SELECT * FROM R1 NATURAL JOIN R2;
```

A	B	C
a ₂	b ₁	c ₁

LEFT OUTER JOIN

- Every tuple in left table must appear in result
- If no matching tuple
 - Padded with NULL values for attributes of right table

```
SELECT * FROM R1 LEFT JOIN R2 ON R1.B = R2.B;
```

R1	
A	B
a ₁	ω
a ₂	b ₁

R2	
B	C
b ₁	c ₁
b ₂	c ₁
b ₃	c ₁

R1 LEFT JOIN R2			
A	B	B	C
a ₁	ω	ω	ω
a ₂	b ₁	b ₁	c ₁

RIGHT OUTER JOIN

- Every tuple in right table must appear in result
- If no matching tuple
 - Padded with NULL values for attributes of right table

```
SELECT * FROM R1 RIGHT JOIN R2 ON R1.B = R2.B;
```

R1

A	B
a ₁	ω
a ₂	b ₁

R2

B	C
b ₁	c ₁
b ₂	c ₁
b ₃	c ₁

R1 RIGHT JOIN R2

A	B	B	C
a ₂	b ₁	b ₁	c ₁
ω	ω	b ₂	c ₁
ω	ω	b ₃	c ₁

FULL OUTER JOIN

- All tuples from both relations must appear in result

```
SELECT * FROM R1 FULL JOIN R2 ON R1.B = R2.B;
```

R1

A	B
a ₁	ω
a ₂	b ₁

R2

B	C
b ₁	c ₁
b ₂	c ₁
b ₃	c ₁

R1 RIGHT JOIN R2

A	B	B	C
a ₁	ω	ω	ω
a ₂	b ₁	b ₁	c ₁
ω	ω	b ₂	c ₁
ω	ω	b ₃	c ₁

Nested Query

- Retrieve first names of students that passed M214

```
SELECT FName
FROM STUDENT s
WHERE s.StudId IN
      (SELECT StudId FROM GRADES
       WHERE CourId = 'M214' AND Grade IS NOT
        NULL) ;
```

FName
James

- A nested query defined by using IN (or =ANY) operator can be expressed as a single block query

```
SELECT FName
FROM STUDENT s, GRADES g
WHERE s.StudId = g.StudId AND g.CourId = 'M214'
AND g.Grade IS NOT NULL;
```

EXISTS and NOT EXISTS

- *Retrieve Id's and surnames of students who passed at least one course:*

```
SELECT s.StudId, s.LName FROM STUDENT s
WHERE EXISTS
    (SELECT * FROM GRADES
     WHERE s.StudId = StudId AND Grade IS NOT NULL) ;
```

- *Retrieve Id's and surnames of students who didn't pass any course:*

```
SELECT s.StudId, s.LName FROM STUDENT s
WHERE NOT EXISTS
    (SELECT * FROM GRADES
     WHERE s.StudId = StudId AND Grade IS NOT NULL ) ;
```

Aggregate Functions

Q2: What is the total point value of the courses in the COMP dept?

```
SELECT SUM(Points)
FROM COURSE
WHERE Dept = 'Comp';
```

sum_Points
52

Q3: How many different first names do COMP majors have?

```
SELECT COUNT(DISTINCT FName) AS NoOfNames
FROM STUDENT
WHERE Major = 'Comp';
```

NoOfNames
1

Grouping

For each student, retrieve the number of courses passed

```
SELECT StudId, COUNT (*)  
FROM GRADES  
WHERE Grade IS NOT NULL  
GROUP BY StudId ;
```

StudId	COUNT(*)
007007	4
131313	1
555555	1

HAVING Clause

Retrieve the number of courses passed for students that passed at least two courses

```
SELECT StudId, COUNT (*)  
FROM STUDENT s NATURAL JOIN GRADES g  
WHERE Grades IS NOT NULL  
GROUP BY s.StudId  
HAVING COUNT (*) > 1;
```

StudId	COUNT(*)
007007	4

SQL Views

- Create a view

```
CREATE VIEW StudOccupied AS
  SELECT g.StudId, SUM(Hours ) AS Occupied
  FROM Grades g, Course p
  WHERE g.CourId = p.CourId AND Grade IS NULL
  GROUP BY StudId ;
```

- Deleting a view

```
DROP VIEW StudOccupied;
```