

The Relational Data Model

Tutorial

SWEN 304
Trimester 2, 2017

Lecturer: Dr Hui Ma

Engineering and Computer Science



Outline

- Relation Schema and its instances
- Key constraints
 - Definitions
 - Procedure of identify keys
- Relational database schema and its instances
- referential integrity constraints
 - Definition
 - Algorithm
- Relational Database operations and constraints

Redefining Some Terms

- **Relation schema** $N(R, C)$
 - N is the name, R is the set of attributes, C is the set of constraints
- Relation schema **instance** $r(N)$:
 - Relation over R that satisfies all constraints from C
- **Tuple**: the set of pairs $t = \{(A_1, a_1), \dots, (A_n, a_n)\}$, where $A_i \in R$, $a_i \in \text{dom}(A_i)$, and $n = |R|$ is $\text{Degree}(r(N))$
- A relation schema instance is a **set of tuples**

Relation Schema Key and Primary Key

- Let $N(A_1, \dots, A_n)$ be a relation schema and $X = \{A_k, \dots, A_m\} \subseteq \{A_1, \dots, A_n\}$, X is a **relation schema key** of N , if

1° $(\forall r(N))(\forall u, v \in r(N))(u[X] = v[X] \Rightarrow u = v)$ (**unique**)

2° $(\forall Y \subset X)(\neg 1^\circ)$ (**minimal**)

3° $(\forall r(N))(\forall t \in r(N))(\forall A \in X)(t[A] \neq \omega)$ (**not null**)

- Relation schema keys are also called **candidate keys** or **keys**
- One of the candidate keys is designated as a **primary key** of the relation

Example keys

- Example
 - CAR (LicPlateNo, EngineNo, Make, Model, Year)
 - Relation schema key: $K = \{\text{LicPlateNo}, \text{EngineNo}\}$,
 - Primary key: $K_p = \{\text{LicPlateNo}\}$
 - CAR (LicPlateNo, EngineNo, Make, Model, Year)
 - the primary key is underlined

Key Constraints

- You are given a relation schema $N(R, C)$ and an instance $r(N)$
- Suppose C does not contain any key specification
- Inferring keys from instances is very **hard** if possible at all, since there are so many of them
- By analyzing instances and $Null(N, A)$ constraints, you can only conclude which subsets of R can **not** be a key
- Also, from instances you may infer which key constraints are **not violated** by instances

Find Key Constraints not Violated in $r(N)$

a) Suppose $Null(N, A) = N$ for all attributes except F in N_2

$r(N_1) =$

A	B	C	D
a_1	b_1	c_1	d_1
a_2	b_2	c_2	d_2
a_3	b_3	c_3	d_3
a_4	b_3	c_4	d_3
a_5	b_1	c_5	d_3

$r(N_2) =$

A	B	C	D	E	F
a_1	b_1	c_1	d_1	e_1	f_1
a_1	b_2	c_1	d_2	e_1	f_2
a_2	b_1	c_2	d_1	e_2	f_3
a_1	b_3	c_3	d_1	e_1	ω
a_3	b_1	c_1	d_3	e_2	f_4

b) Suppose now $Null(N_1, C) = Y$ and $Null(N_2, D) = Y$ and $Null(N_2, F) = Y$, and there are some null values in the corresponding columns

Procedure

1. Produce a **power** set of the set of relation schema attributes
2. Check no-empty subsets for key constraint satisfaction, starting from the subsets with **lower** cardinality
3. If a subset satisfies a key constraint, all its **supersets** will also satisfy, and therefore do not need to be checked
 - Results ($SatKey(N)(r(N))$) key constraint of relation schema N not violated in $r(N)$
 - a):
 - $SatKey(N_1)(r(N_1)) = \{A, C\},$
 - $SatKey(N_2)(r(N_2)) = \{AB, CD, BCE, BDE\}$
 - b):
 - $SatKey(N_1)(r(N_1)) = \{A\},$
 - $SatKey(N_2)(r(N_2)) = \{AB, BCE\}$

Relational Database Schema

- **Relational database schema** $N(S, IC)$
 - N is the name,
 - $S = \{N_1(R_1, C_1), \dots, N_k(R_k, C_k)\}$ is a set of relation schemas, and
 - IC is a set of **interrelation** constraints
- A database schema DBS as a **complex** data type defines a finite, but very large number of different **database instances**
- An instance of the relational database schema $N(S, IC)$ is $db = \{r(N_1), \dots, r(N_k)\}$ such that:
 - Each $r(N)$ is an instance of a relation schema $N(R, C)$ in S , and
 - db satisfies all constraints in IC

Referential Integrity

- Given $N_1(R_1, C_1)$ and $N_2(R_2, C_2)$ with X the primary key of N_1 and $Y = \{B_1, \dots, B_m\} \subseteq R_2$, Y is a foreign key in N_2 with regard to X in N_1

- Relations $r(N_1)$ and $r(N_2)$ satisfy the **referential integrity constraint** $N_2[Y] \subseteq N_1[X]$ if:

$$(\forall u \in r(N_2))(\exists v \in r(N_1))(u[Y] = v[X] \vee (\exists i \in \{1, \dots, m\})(u[B_i] = \omega))$$

- N_2 : the referencing relation schema and
 - N_1 : the referenced relation schema
- Either tuples u and v are **equal** on X and Y values, or there exists at least one attribute in Y whose value in the tuple u **is null**

Define Referential Integrity Constraints

- Suppose each set S of relation schemas is designed using a “disciplined” approach that subsumes at least the following:
 - Of all keys of a relation schema, only the primary key attributes can be propagated to the other relation schemes,
 - Hence, relation schemes in S are interconnected by (foreign key, primary key) pairs, (further (F, P) pair)
 - When a relation schema N_2 contains a foreign key, and a relation schema N_1 contains the primary key of a (F, P) pair and the attribute A belongs to the both primary and foreign keys, then:
 - $Dom(N_2, A) \subseteq Dom(N_1, A)$, and
 - $Range(N_2, A) \subseteq Range(N_1, A)$
 - It follows that attributes $N_1.A$ and $N_2.A$ have the same meaning (why should they otherwise have the same name and domain?)

Algorithm

- Input: $S = \{N_i(R_i, C_i) \mid i = 1, 2, \dots, n\}$
- Output: $IC = \{N_i[X] \subseteq N_j[X] \mid N_i \in S, N_j \in S, i, j \in \{1, 2, \dots, n, i \neq j\}\}$

- Procedure:

Set $IC = \{ \}$

For each pair relation schemas $N_1(R_1, C_1)$ and $N_2(R_2, C_2)$ in S

If $R_1 \cap R_2 \supseteq X$ and X is the primary key of (say) N_1 , then

$IC = IC \cup \{N_2[X] \subseteq N_1[X]\},$

For each $ic : N_2[X] \subseteq N_1[X] \in IC$

If $IC \models ic : N_2[X] \subseteq N_1[X]$, then

Set $IC = IC \setminus \{ic\}$

(\models logically implies)

Exception

- The algorithm for defining referential integrity constraints will not discover those referential integrity constraints that involve attribute Y in N_2 and attribute X in N_1 , which have different names, but the same semantic (which are also domain compatible)
- To resolve these situations we need the information of the form:
 - $Dom(N_2, Y) \subseteq Dom(N_1, X)$,
 - $Range(N_2, Y) \subseteq Range(N_1, X)$, and (consequently)
 - (N_2, Y) and (N_1, X) have the same semantics

Incorrect Referential Integrity Constraints

$S = \{ \text{BOOK}(\underline{\text{ISBN}}, \text{Title}), \text{LIBRARY}(\underline{\text{LibId}}, \text{LibN}),$
 $\text{BOOK_COPIES}(\underline{\text{ISBN}}, \underline{\text{LibId}}, \text{CopNum}),$
 $\text{BOOK_LOANS}(\underline{\text{ISBN}}, \underline{\text{LibId}}, \underline{\text{CardNo}}, \text{Date}),$
 $\text{BORROWER}(\underline{\text{CardNo}}, \text{Name}) \}$

$IC = \{ \text{BOOK_COPIES}[\text{ISBN}] \subseteq \text{BOOK}[\text{ISBN}],$
 $\text{BOOK_COPIES}[\text{LibId}] \subseteq \text{LIBRARY}[\text{LibId}],$
 $\text{BOOK_LOANS}[\text{CardNo}] \subseteq \text{BORROWER}[\text{CardNo}],$
 $\text{BOOK_LOANS}[\text{ISBN}] \subseteq \text{BOOK}[\text{ISBN}],$
 $\text{BOOK_LOANS}[\text{LibId}] \subseteq \text{LIBRARY}[\text{LibId}],$
 $\text{BOOK_LOANS}[\text{ISBN}] \subseteq \text{BOOK_COPIES}[\text{ISBN}],$
 $\text{BOOK_LOANS}[\text{LibId}] \subseteq \text{BOOK_COPIES}[\text{LibId}]$
 $\}$

Are the constraints correct?

A Consequence of Incorrect Instance

BOOK

ISBN	Title
1010	DB Sys
9999	Comp

LIBRARY

LibId	LibN
1	Vic
9	Massey

BORROWER

CardNo	Name
10	Susan
20	James

BOOK_COPIES

ISBN	LibId	NoOfCop
1010	1	10
9999	1	15
9999	9	5

BOOK_LOANS

ISBN	LibId	CardNo	Date
1010	1	10	01.03.01
9999	1	10	15.07.00
1010	9	20	01.03.01

$IC = \{ \text{BOOK_COPIES}[\text{ISBN}] \subseteq \text{BOOK}[\text{ISBN}],$
 $\text{BOOK_COPIES}[\text{LibId}] \subseteq \text{LIBRARY}[\text{LibId}],$
 $\text{BOOK_LOANS}[\text{CardNo}] \subseteq \text{BORROWER}[\text{CardNo}],$
 $\text{BOOK_LOANS}[\text{ISBN}] \subseteq \text{BOOK}[\text{ISBN}],$
 $\text{BOOK_LOANS}[\text{LibId}] \subseteq \text{LIBRARY}[\text{LibId}],$
 $\text{BOOK_LOANS}[\text{ISBN}] \subseteq \text{BOOK_COPIES}[\text{ISBN}],$
 $\text{BOOK_LOANS}[\text{LibId}] \subseteq \text{BOOK_COPIES}[\text{LibId}] \}$

A Consequence of Incorrect RI

BOOK

ISBN	Title
1010	DB Sys
9999	Comp

LIBRARY

LibId	LibN
1	Vic
9	Massey

BORROWER

CardNo	Name
10	Susan
20	James

BOOK_COPIES

ISBN	LibId	NoOfCop
1010	1	10
9999	1	15
9999	9	5

BOOK_LOANS

ISBN	LibId	CardNo	Date
1010	1	10	01.03.01
9999	1	10	15.07.00
1010	9	20	01.03.01

Massey library doesn't possess the book DB Sys

Wrong tuple

Incorrect Referential Integrity Constraints

$S = \{ \text{BOOK}(\underline{\text{ISBN}}, \text{Title}), \text{LIBRARY}(\underline{\text{LibId}}, \text{LibN}),$
 $\text{BOOK_COPIES}(\underline{\text{ISBN}}, \underline{\text{LibId}}, \text{CopNum}),$
 $\text{BOOK_LOANS}(\underline{\text{ISBN}}, \underline{\text{LibId}}, \underline{\text{CardNo}}, \text{Date}),$
 $\text{BORROWER}(\underline{\text{CardNo}}, \text{Name}) \}$

$IC = \{ \text{BOOK_COPIES}[\text{ISBN}] \subseteq \text{BOOK}[\text{ISBN}],$
 $\text{BOOK_COPIES}[\text{LibId}] \subseteq \text{LIBRARY}[\text{LibId}],$
 $\text{BOOK_LOANS}[\text{CardNo}] \subseteq \text{BORROWER}[\text{CardNo}],$
 $\text{BOOK_LOANS}[\text{ISBN}] \subseteq \text{BOOK}[\text{ISBN}],$
 $\text{BOOK_LOANS}[\text{LibId}] \subseteq \text{LIBRARY}[\text{LibId}],$
 $\text{BOOK_LOANS}[\text{ISBN}] \subseteq \text{BOOK_COPIES}[\text{ISBN}],$
 $\text{BOOK_LOANS}[\text{LibId}] \subseteq \text{BOOK_COPIES}[\text{LibId}],$
 $\text{BOOK_LOANS}[(\text{ISBN}, \text{LibId})] \subseteq \text{BOOK_COPIES}[(\text{ISBN}, \text{LibId})]$
 $\}$

redundant

wrong

missing

Inferring Referential Integrities

- Convince yourself (by thinking) that the following implication is true

$$\begin{aligned}
 &(\text{BOOK_COPIES} [\text{ISBN}] \subseteq \text{BOOK} [\text{ISBN}] \wedge \\
 &\text{BOOK_COPIES} [\text{LibId}] \subseteq \text{LIBRARY} [\text{LibId}] \wedge \\
 &\text{BOOK_LOANS} [(\text{ISBN}, \text{LibId})] \subseteq \text{BOOK_COPIES} [(\text{ISBN}, \text{LibId})]) \\
 &|= \\
 &(\text{BOOK_LOANS} [\text{ISBN}] \subseteq \text{BOOK} [\text{ISBN}] \wedge \\
 &\text{BOOK_LOANS} [\text{LibId}] \subseteq \text{LIBRARY} [\text{LibId}])
 \end{aligned}$$

- But also note that:

$$\begin{aligned}
 &\neg((\text{BOOK_COPIES} [\text{ISBN}] \subseteq \text{BOOK} [\text{ISBN}] \wedge \\
 &\text{BOOK_COPIES} [\text{LibId}] \subseteq \text{LIBRARY} [\text{LibId}] \wedge \\
 &\text{BOOK_LOANS} [\text{ISBN}] \subseteq \text{BOOK} [\text{ISBN}] \wedge \\
 &\text{BOOK_LOANS} [\text{LibId}] \subseteq \text{LIBRARY} [\text{LibId}]) \mid= \\
 &(\text{BOOK_LOANS} [(\text{ISBN}, \text{LibId})] \subseteq \text{BOOK_COPIES} [(\text{ISBN}, \text{LibId})])
 \end{aligned}$$

Correct Referential Integrity Constraints

$$S = \{ \text{BOOK } (\underline{\text{ISBN}}, \text{Title}), \text{LIBRARY } (\underline{\text{LibId}}, \text{LibN}), \\ \text{BOOK_COPIES } (\underline{\text{ISBN}}, \underline{\text{LibId}}, \text{NoOfCop}), \\ \text{BOOK_LOANS}(\underline{\text{ISBN}}, \underline{\text{LibId}}, \underline{\text{CardNo}}, \text{Date}), \\ \text{BORROWER } (\underline{\text{CardNo}}, \text{Name}) \}$$
$$IC = \{ \text{BOOK_COPIES } [\text{ISBN}] \subseteq \text{BOOK } [\text{ISBN}], \\ \text{BOOK_COPIES } [\text{LibId}] \subseteq \text{LIBRARY } [\text{LibId}], \\ \text{BOOK_LOANS } [(\text{ISBN}, \text{LibId})] \subseteq \\ \text{BOOK_COPIES } [(\text{ISBN}, \text{LibId})], \\ \text{BOOK_LOANS } [\text{CardNo}] \subseteq \text{BORROWER } [\text{CardNo}] \}$$

Renaming Attributes with Different Roles (H)

- The referential integrity

$\text{REQ_BOOK}[(\text{ISBN}, \text{LibId})] \subseteq \text{BOOK_COPIES}[(\text{ISBN}, \text{LibId})]$

is incorrect:

- The attributes REQ_BOOK.ISBN and BOOK_COPIES.ISBN have different meanings
- For a given LibId value, REQ_BOOK.ISBN and BOOK_COPIES.ISBN have disjoint sets of values
 - REQ_BOOK.ISBN are ISBNs of books not yet in the library
 - BOOK_COPIES.ISBN are ISBNs of books already in the library
- Instead we use the referential integrity constraints
$$\text{REQ_BOOK}[\text{ISBN}] \subseteq \text{BOOK}[\text{ISBN}]$$
$$\text{REQ_BOOK}[\text{LibId}] \subseteq \text{LIBRARY}[\text{LibId}]$$

to ensure that new books to be purchased are first recorded in the BOOK table (for bookkeeping) and are requested for existing libraries only

Improving Extended Library Schema (H)

- After the correction we have the relation schema
 $\text{REQ_BOOK}(\{\text{CardNo}, \text{ISBN}, \text{LibId}, \text{ReqDate}\},$
 $\{\text{CardNo} + \text{ISBN} + \text{LibId}\})$
- ... and the referential integrity constraints:
 $\text{REQ_BOOK}[\text{ISBN}] \subseteq \text{BOOK}[\text{ISBN}]$
 $\text{REQ_BOOK}[\text{LibId}] \subseteq \text{Library}[\text{LibId}]$
 $\text{REQ_BOOK}[\text{CardNo}] \subseteq \text{BORROWER}[\text{CardNo}]$

Relational Database Operations

- Database Management System must implement **update** operations:
 - **insert**,
 - **delete**, and
 - **modify**

- Database Management System must implement **retrieval** operations:
 - query language
 - Need a well defined language

DB Updates and Constraints

- No update operation should leave a database in an inconsistent state (with violated constraints)
- A DBMS must take the actions necessary to prevent a constraint violation:
 - **reject**: do not allow the operation
 - **cascade**: propagate the operation by making necessary consequential changes
 - **set null**, or **set default**: reset other values to maintain consistency

Inserts and Constraint Violations

- **Inserting** a new tuple could **violate**
 - Attribute/domain constraints
(a value is not of the right type or within the required range)
 - Uniqueness constraints
(the values of the key attributes duplicate another tuple)
 - Not Null constraints
(an attribute has the value null when it shouldn't)
 - Referential Integrity constraints
(the values of the attributes of a foreign key do not match any tuple in the other relation)
- **Response:**
 - **Reject** the operation – there is no change that the DBMS system could safely make to resolve the inconsistency

Deletes and Constraint Violations

- **Deleting** a tuple can only **violate** a **referential integrity constraint**:
 - If a tuple t is referred to by foreign keys in some tuples t_1, t_2, \dots, t_n in other relations, then deleting t will make t_1, t_2, \dots, t_n inconsistent.
 - Example:
 - Delete a student record from the database, and all their grade records will refer to nothing
- There are several options:
 - **Reject** the deletion
 - **Set null / set default**: insert null or a default value in the *foreign key* attributes of tuples in other relation(s) that refer to t (can't do set null if foreign key attributes are NOT NULL)
 - **Cascade**: delete tuples in other relation(s) that refer to t (appropriate only if the other tuples "existentially depend" on t)

Modify and Constraint Violations

- **Modifying/updating** the values of attributes in a tuple may **violate** constraints
 - Attribute/domain constraints
Response: **reject** (like insert)
 - Key constraints (if attribute is part of a key)
Response: treat as a **delete followed by an insert**
 - Referential integrity constraints (if attribute is part of a foreign key).
Response: **reject** (like insert), or **cascade**, or **set null**, or **set default** (like delete)

DB Updates and Constraints

Update operation	Domain / Attribute constraint	Key / Entity integrity constraint,	Referential integrity
insert	reject	reject	reject
delete	no violation	no violation	reject, cascade, set null, set default
modify	reject	reject	reject, cascade, set null, set default

A Question for You

- Consider the following database instance

TEXTBOOK			
Title	<u>ISBN</u>	Pcod	Pnum
COD	1111	COMP	203
FDBS	2222	COMP	ω

COURSE		
<u>Pcode</u>	<u>Pnum</u>	Pname
COMP	203	CO
COMP	302	DBS

- Should a DBMS reject the following update operation:
(Y/N)?

```
UPDATE TEXTBOOK SET PNum = 302 WHERE ISBN = 2222;
```

N

- Should a DBMS reject the following update operation:
(Y/N)?

```
UPDATE TEXTBOOK SET PNum = 302 WHERE ISBN = 1111;
```

N

A Question for You

- Consider the following database instance

TEXTBOOK			
Title	<u>ISBN</u>	Pcod	Pnum
COD	1111	COMP	203
FDBS	2222	COMP	ω

COURSE		
<u>Pcode</u>	<u>Pnum</u>	Pname
COMP	203	CO
COMP	302	DBS

- Should a DBMS reject the following update operation:
(Y/N)?

```
UPDATE TEXTBOOK SET PNum = 403 WHERE ISBN = 2222;
```

Y

- Should a DBMS reject the following update operation:
(Y/N)?

```
UPDATE COURSE SET PNum = 102 WHERE Pname = 'CO';
```

Y/N

A Question for You

- Consider the following database instance

TEXTBOOK			
Title	<u>ISBN</u>	Pcod	Pnum
COD	1111	COMP	203
FDBS	2222	COMP	ω

COURSE		
<u>Pcode</u>	<u>Pnum</u>	Pname
COMP	203	CO
COMP	302	DBS

- Should a DBMS reject the following update operation:
(Y/N)?

```
UPDATE TEXTBOOK SET Pcode = 'SWEN' WHERE ISBN = 2222;
```

N