

Recovery

SWEN 304

Trimester 2, 2017

Lecturer: Dr Hui Ma

Engineering and Computer Science



Slides by: Pavle Morgan & Hui Ma

Plan For The Recovery Topic

- Transaction Log File
- Classification of database recovery procedures
 - Deferred database update
 - Immediate database update
- Checkpoints
 - *Readings:*
 - *Chapter 23 of the textbook*

Purpose of Database Recovery

- To bring the database into the last consistent state, which existed prior to the failure
- To preserve transaction properties (**A**tomicity, **C**onsistency, **I**solation and **D**urability)
- Example:
 - If the system crashes before a fund transfer transaction completes its execution, then either one or both accounts may have incorrect value.
 - Thus, the database must be restored to the state before the transaction modified any of the accounts

Transaction Log

- To be able to recover from failures DBMS maintains a log file
- For recovery from any type of failure data values prior to modification (BFIM - BeFore IMage) and the new value after modification (AFIM – AFTer IMage) are required
- Typically, a log file contains records with the following contents:

[start_transaction, T] (* T is a transaction id*)

[write_item, T , X , old_value, new_value]

[read_item, T , X] (*optional*)

[commit, T]

[abort, T]

Transaction Roll-back (Undo) and Roll-Forward (Redo)

- To maintain atomicity, a transaction's operations are **redone** or **undone**
 - **Undo**: Restore all BFIMs on to disk (Remove all AFIMs)
 - **Redo**: Restore all AFIMs on to disk
- Database recovery is achieved either by performing only Undos or only Redos or by a combination of the two
- These operations are recorded in the log as they happen

Classification of database recovery procedures

- According to the type of a failure, recovery procedures are classified as:
 - Recovery from a **catastrophic** failure (like disk crash), and
 - Recovery from a **noncatastrophic** failure
- Recovery from a **catastrophic** failure is based on restoring a database **back_up** copy by redoing operations of committed transactions (stored in archived **log** files) up to the time of the failure

Noncatastrophic Failures

- A computer failure (system crash):
 - E.g. a hardware, software or network error occurs in the computer system
- A transaction or system error:
 - E.g. integer overflow, division by zero, logical error, user interruption
- Local errors or exception conditions detected by transactions
 - E.g. Data not found, exception condition
- Concurrency control enforcement:
 - E.g. violate serializability, deadlock

Classification (continued)

- If a database becomes inconsistent due to a **noncatastrophic** failure, the strategy is to reverse only those changes that made database inconsistent
- It is accomplished by **undoing** (and sometimes also **redoing**) some operations, with the use of an in memory log file
- From now on we consider only recovery from non disk crash failures (we suppose data on disk are safe)
- The recovery from noncatastrophic failures can be based on many algorithms, as:
 - **Deferred** update,
 - **Immediate** update, and
 - Shadow update (not discussed)

Database Recovery

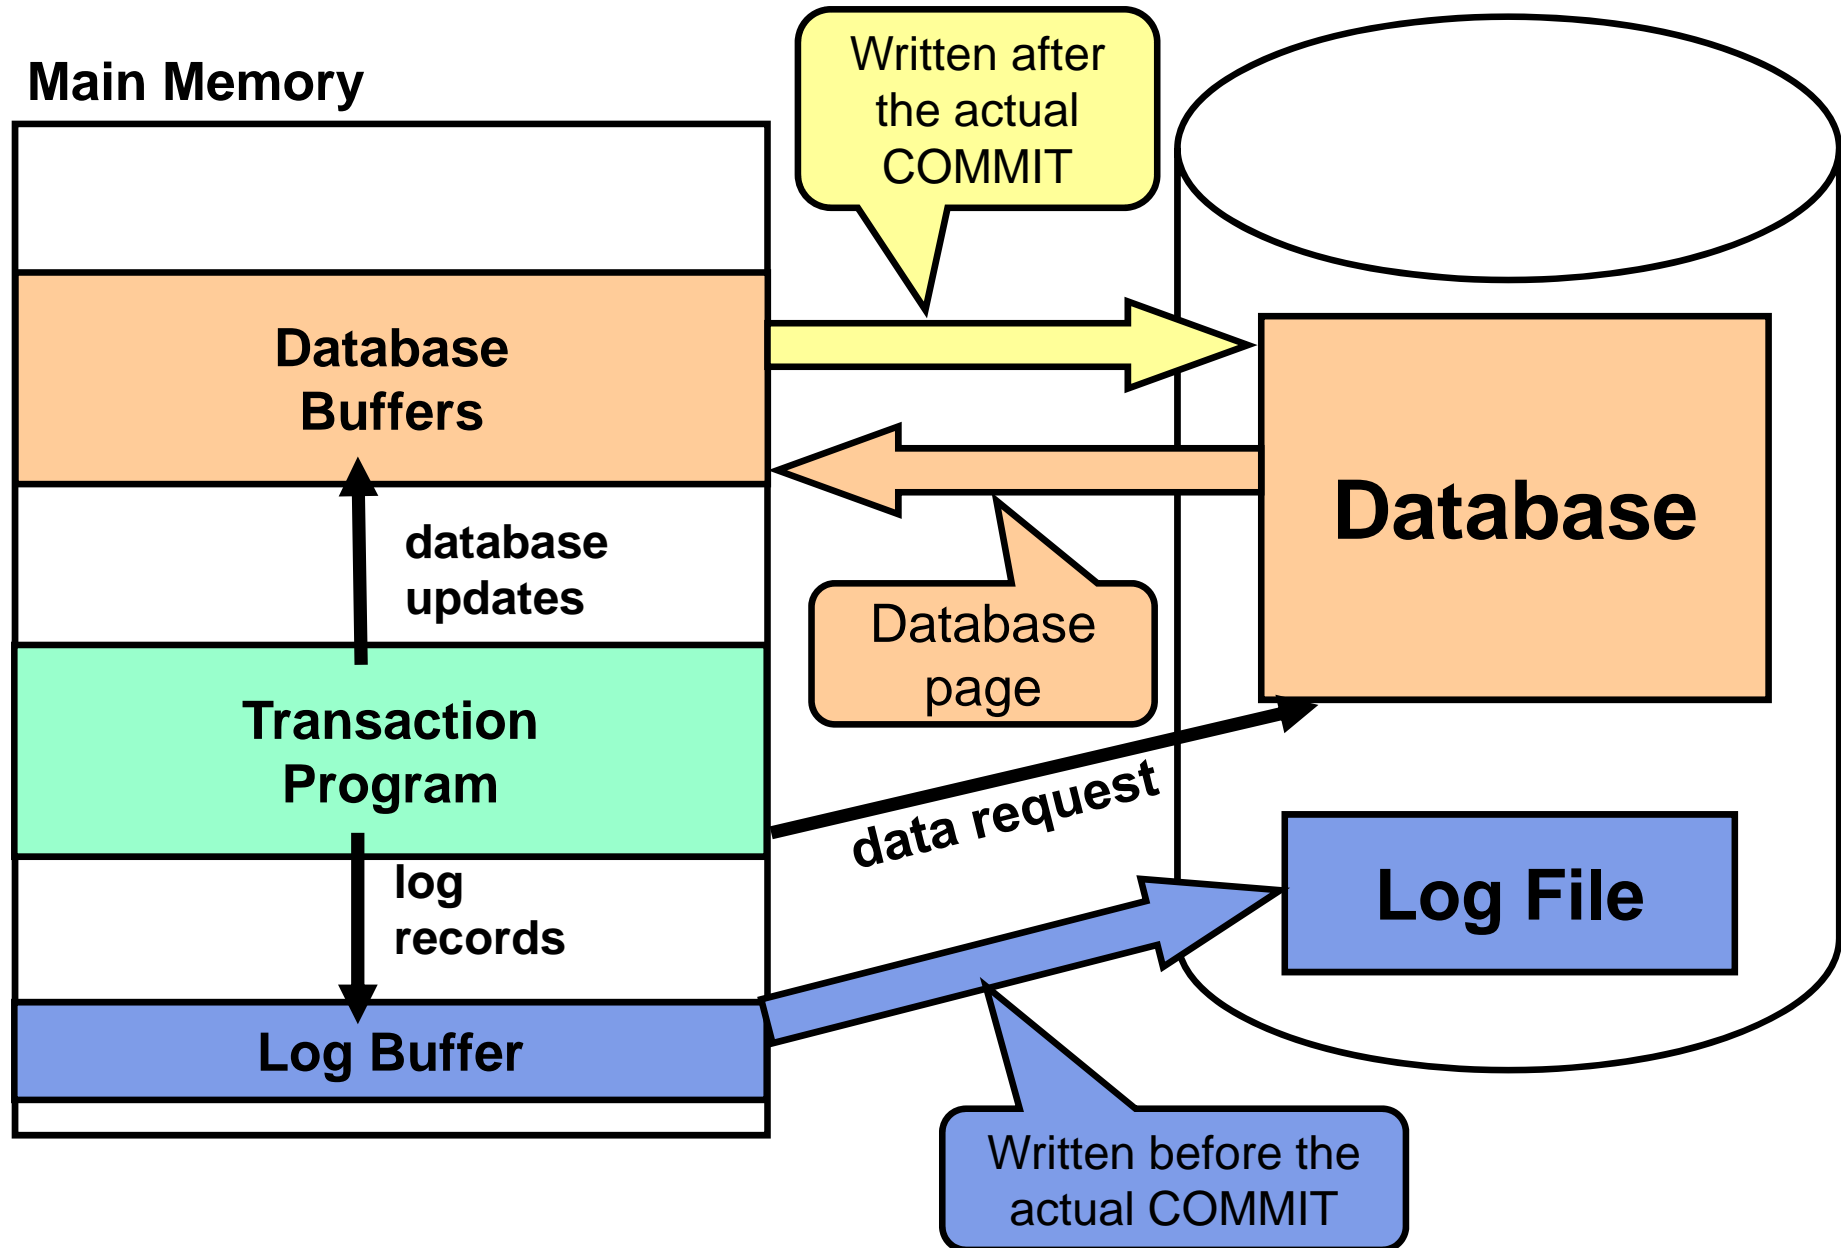
■ Data Update

- **Deferred Update:** All modified data items in the cache is written either after a transaction ends its execution or after a fixed number of transactions have completed their execution
- **Immediate Update:** As soon as a data item is modified in cache, the disk copy is updated
- **Shadow update:** The modified version of a data item does not overwrite its disk copy but is written at a separate disk location

Deferred Update

- The idea:
 - Postpone updates to the database until a transaction reaches its **commit** command
- Updates are recorded in a log file and in cache buffers with database pages (all in RAM)
- When the COMMIT is reached, before it is executed all log updates are first force written to the **log file** on disk, and then the transaction commits
- After that, corresponding updates are written from **buffers** to the database

Deferred Update Layout



Deferred Update (continued)

- If a transaction fails before reaching COMMIT, there is no need to make any recovery
- If a system crash occurs after COMMIT, but before all changes are recorded in the database on disk, **Redo** of operations is needed,
 - The operations have to be redone from the log file (that is already on disk) to the database
 - Using after images, AFIMs (item values intended to be written to the database) to perform redo
 - Deferred update recovery log file has to contain only after images - the **new** database item values

Deferred Update (An Example)

Operations of T_1	Generated log records
start_ T_1	begin, $\langle T_1 \rangle$
read_item(A) // $A = 4000$ $A = A - 1000$	
write_item(A)	$\langle T_1, \text{write_item}(A), 3000 \rangle$
read_item(B) // $B = 0$ $B = B + 1000$	
write_item(B)	$\langle T_1, \text{write_item}(B), 1000 \rangle$
read_item(A) // $A = 3000$ $A = A - 1000$	
write_item(A)	$\langle T_1, \text{write_item}(A), 2000 \rangle$
read_item(C) // $C = 0$ $C = C + 1000$	
write_item(C)	$\langle T_1, \text{write_item}(C), 1000 \rangle$
commit T_1	
	//force write to the log on disk
	commit, $\langle T_1 \rangle$ // actual


Deferred Update

- To finish a transaction, DBMS has force stored log records on disk
- If the system fails after that point, DBMS will use the log records to REDO changes in the database
- It is sufficient to redo only the **last** written value (after image) of every item changed
- So, redoing starts from the end of the log file and maintains a list of already redone database items

Checkpoints

- Some DBMS use **CHECKPOINT** records in the log file to prevent unnecessary redo operations
- To take a checkpoint record, DBMS has:
 - To suspend temporarily operations of all transactions,
 - To force write results of all update operations of committed transactions from main memory buffers to disk,
 - Write a checkpoint record into the log file and write log to disk
 - Resume executing transactions
- Only changes made by transactions committed between the last checkpoint and a system failure have to be redone

Redo With Checkpoint



Changes made by T_1 are stored in the database, so there is nothing to do with them

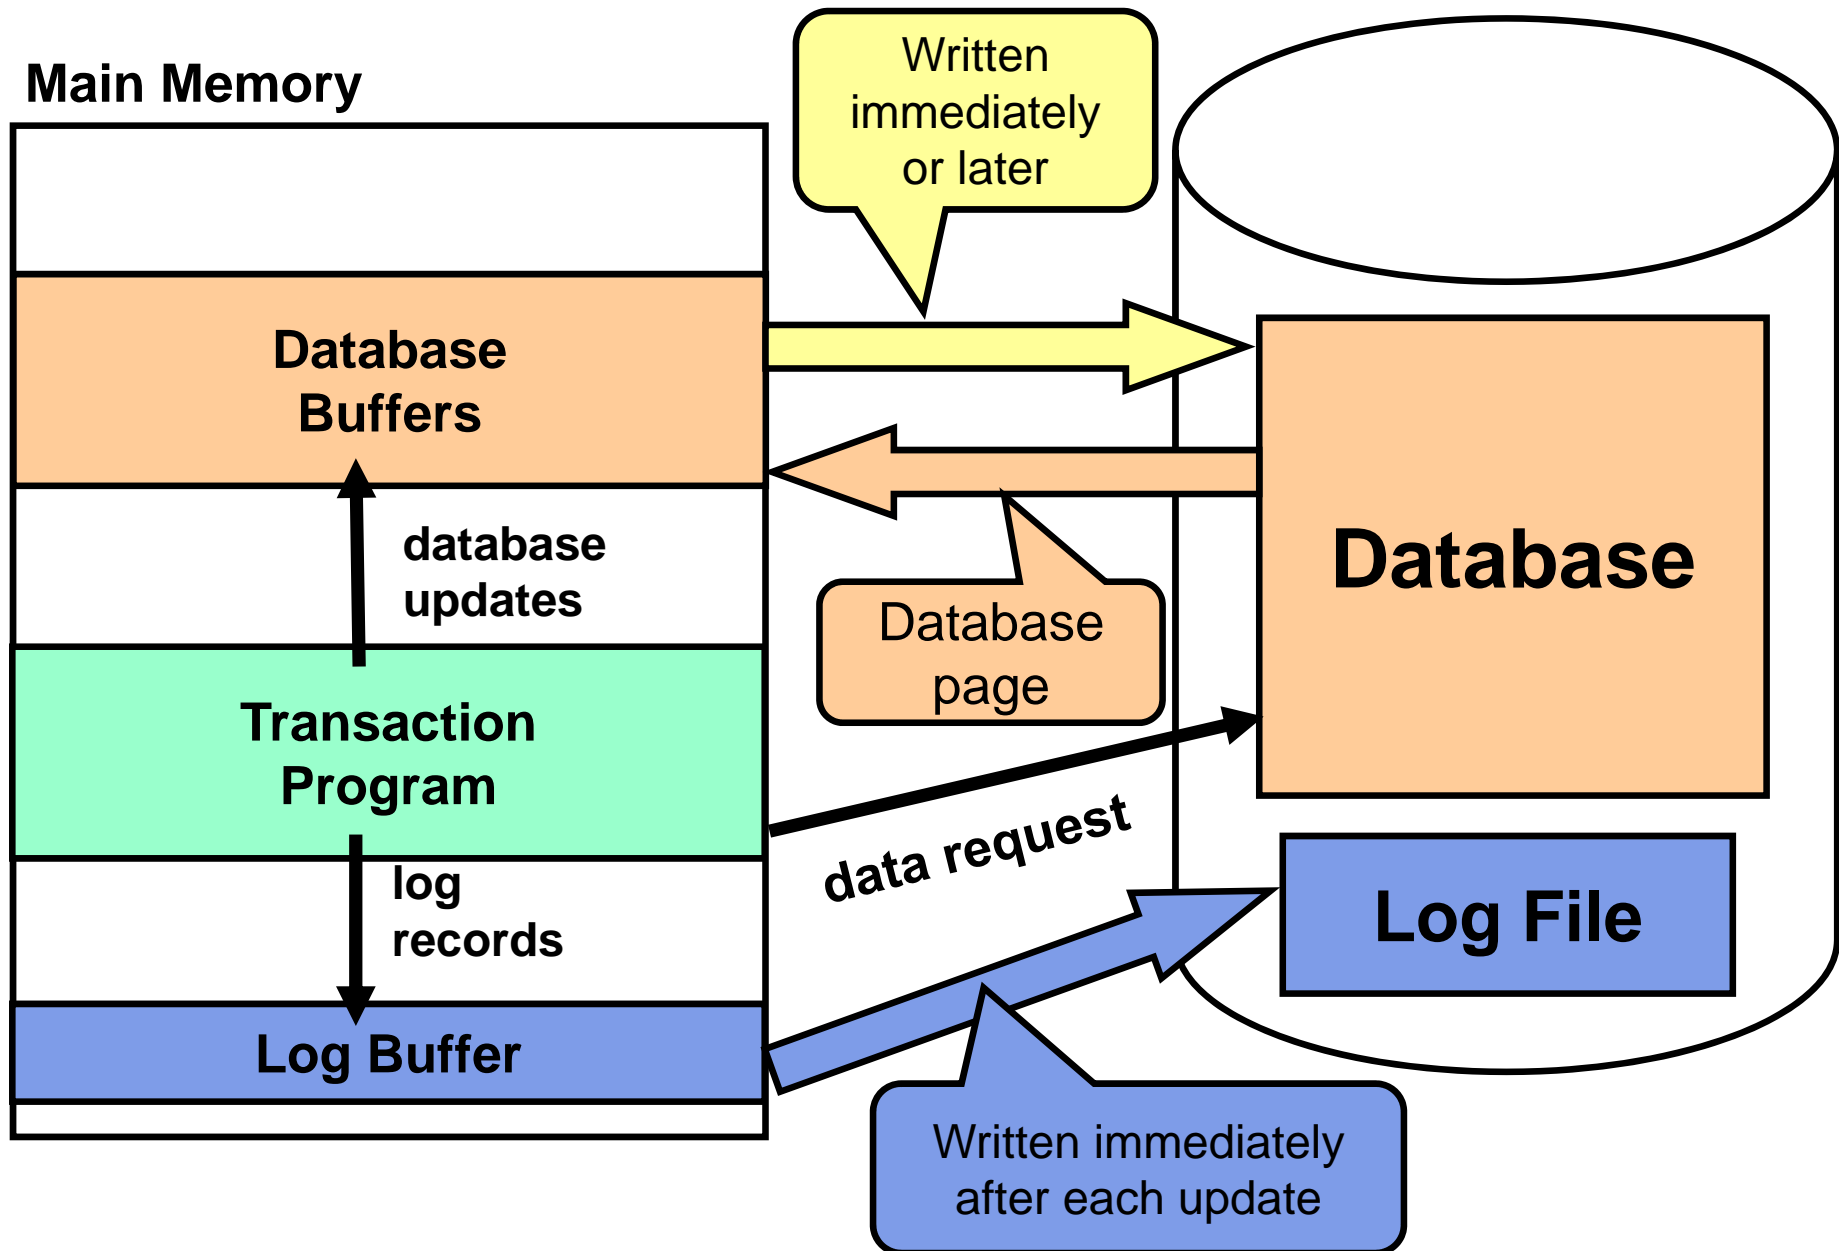
DBMS will **REDO** transactions T_2 and T_3 (their logs are already on disk)

DBMS or the user has to **rerun** T_4 and T_5

Immediate Update

- The main idea:
 - When a transaction issues an update command, before and after images are recorded into a log file on **disk**, and (thereupon) the database **can** (but do not have to) be **immediately** updated
- There are two versions of immediate update algorithm
 - One writes all updates from buffers into the database **before** the transaction actually **commits**
 - The other one allows a transaction to **commit before** all its updates have been written in the database
 - The second version is the most common one

Immediate Update Layout



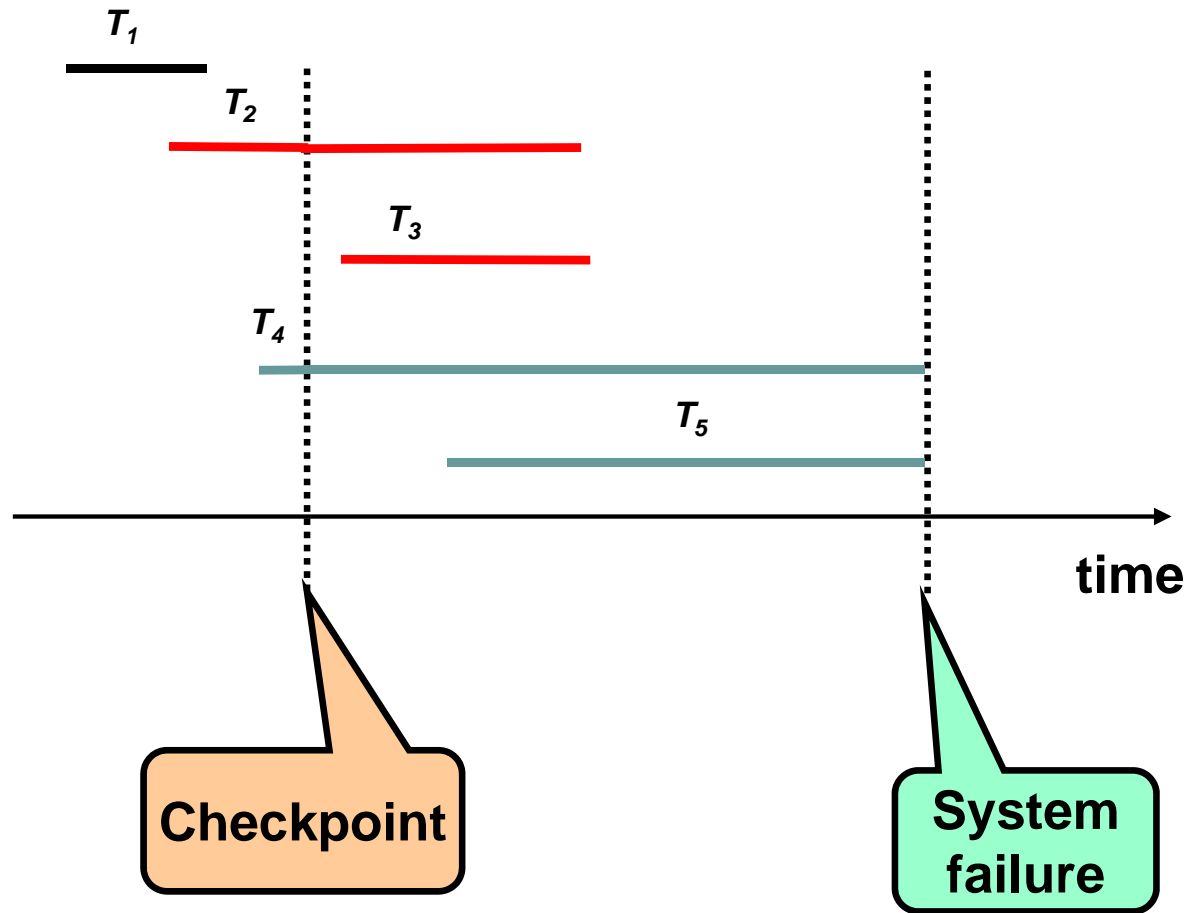
Transaction Roll-Back (Undo)

- If a transaction fails for any reason, its effects are rolled-back
- **Roll-back** of a transaction is a procedure of **undoing** changes made against the database by a non-committed transaction
- A roll-back is done by restoring the **before images** (old values) of each database item changed by the transaction
- The roll-back is done in the **reverse order** of the order the operations were written into the log file on disk

Redo in Immediate Update

- In the case of a system crash:
 - Because there can be some updates of the committed transactions that are not written into the database, effects of the committed transactions are redone to the database
 - Recall, we allow a transaction to commit before all of its updates have been written into the database
- Only the **last** updates of each item have to be redone

Immediate Update



Changes made by T_1 are stored in the database, so there is nothing to do with them

DBMS will **REDO** transactions T_2 and T_3 (their logs are already on disk)

DBMS has to **UNDO** T_4 and T_5

A Question for You

- Immediate update with roll-back requires redoing transactions that committed between the last checkpoint and the moment of system crash
- What if the changes are already made against the database?
 - a) Redoing will bring the database into an inconsistent state
 - b) This question is too complex, and can not be answered without a careful analysis
 - c) Redoing will make no harm to the database, because it will bring the database in the same consistent state as it was in

Summary

- *Recovery* from a *catastrophic* failure is made by applying operations of committed transactions from archived log files on an archived database back up copy
- Recovery from a *noncatastrophic* failure can be accomplished through:
 - Deferred update,
 - Immediate update, and
 - Other recovery schemes
- *Deferred update* means writing changes into a database after a transaction commits, but the log is written to disk just before the transaction actually commits
- *Immediate update* means that changes are immediately stored in a log file on disk. So, changes can be written into a database before or after a transaction commits